

PLEASE CLICK 'DOWNLOAD' – THERE ARE OVER 200 SLIDES AND GITHUB WILL NOT DISPLAY ALL OF THEM

Artificial Intelligence and the Future of Medicine/Mankind

Howard Schneider
Sheppard Clinic North, Toronto, Canada

Ontario Medical Association (OMA)
Section on Primary Care Mental Health
Toronto, October 22, 2020

Artificial Intelligence and the Future of Medicine/Mankind

Howard Schneider
Sheppard Clinic North, Toronto, Canada

Ontario Medical Association (OMA)
Section on Primary Care Mental Health
Toronto, October 22, 2020

NEW TOPIC

PRELIMINARIES....

3

NO MEDICAL COMMERCIAL NAMES

- ▶ but....many non-medical AI-related products are company specific

CONFLICTS OF INTEREST

- ▶ None

DR HOWARD SCHNEIDER, MD, MDPAC(C), CCFP, FCFP

WHY SHOULD I GIVE THIS TALK?

Practicing physician

- ▶ Background –General Practice/ Family Practice
- ▶ A decade –Psychiatry Consultations in ER (Laval)
- ▶ Two decades –Community Mental HealthCare(GTA)

DR HOWARD SCHNEIDER, MD, MDPAC(C), CCFP, FCFP

WHY SHOULD I GIVE THIS TALK?

Interest – How does mind work?

Flip side – How to create an AGI?

(==*Artificial General Intelligence*, ‘HLAI’, ‘Strong AI’)

- ▶ Research – Cognitive Architectures
- ▶ (I design machines to produce AI/AGI)

LEARNING OBJECTIVES

- ▶ **1. Real understanding of what AI is:**
 - ▶ 1a. Deep Learning and Reinforcement Learning
 - ▶ 1b. Field of Artificial Intelligence (AI)
 - ▶ 1c. Neuro-Symbolic Gap
- ▶ **2. How will AI in next decade (or two) affect my patients' lives?**
- ▶ **3. How will AI affect my practice of medicine including psychotherapy?**
 - ▶ 3a. How is AI affecting medicine at present?
 - ▶ 3b. How will AI affect medicine in the next decade?
- ▶ **4. How will AI affect the future of mankind?**
- ▶ **5. Discussion**

WILL DO BEST TO KEEP ON SCHEDULE....

(90 MINUTES TALK, 30 MINUTES QUESTIONS)



NOT REQUIRED.... BUT.....
CONSIDER MAKING NOTES IF THIS HELPS
YOU LEARN....



To download a copy of these slides:

SEARCH (EG, GOOGLE):

“DR HOWARD SCHNEIDER TORONTO”

OR ENTER:

[HTTPS://GITHUB.COM/HOWARD8888/AI-AND-THE-FUTURE-OF-MEDICINE/BLOB/MASTER/AI%20AND%20THE%20FUTURE%20OF%20MEDICINE.PDF](https://github.com/howard8888/AI-AND-THE-FUTURE-OF-MEDICINE/blob/master/AI%20AND%20THE%20FUTURE%20OF%20MEDICINE.PDF)



LEARNING OBJECTIVES

- ▶ **1. Real understanding of what AI is:**
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ **2. How will AI in next decade (or two) affect my patients' lives?**
- ▶ **3. How will AI affect my practice of medicine including psychotherapy?**
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ **4. How will AI affect the future of mankind?**
- ▶ **5. Discussion**


NEW TOPIC

INTRODUCTION

13

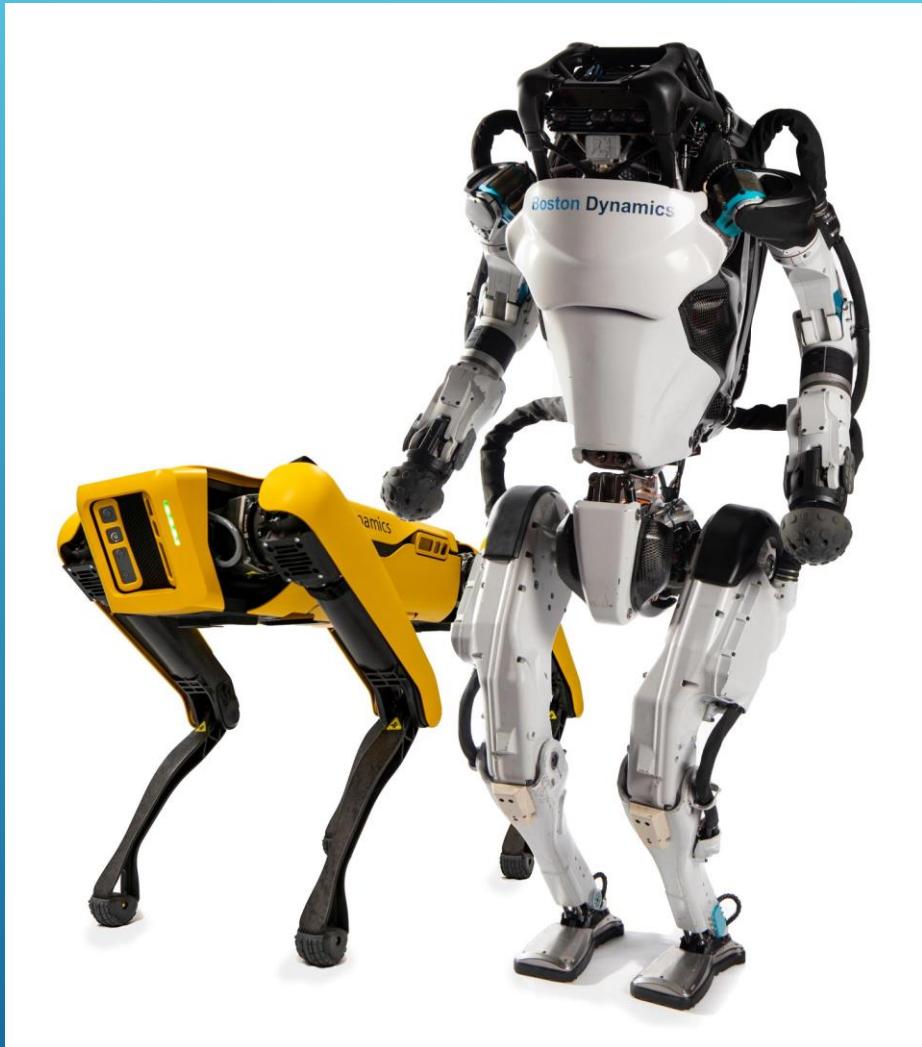
October 2020

FORGET ABOUT CURRENT HYPE YOU SEE
ABOUT AI (ARTIFICIAL INTELLIGENCE) ALL
AROUND YOU IN THE MEDIA....



14

October 2020



-Human-like (and
dog-like) robots
do exist

-They **can** walk,
run, jump, even
do back-flips....

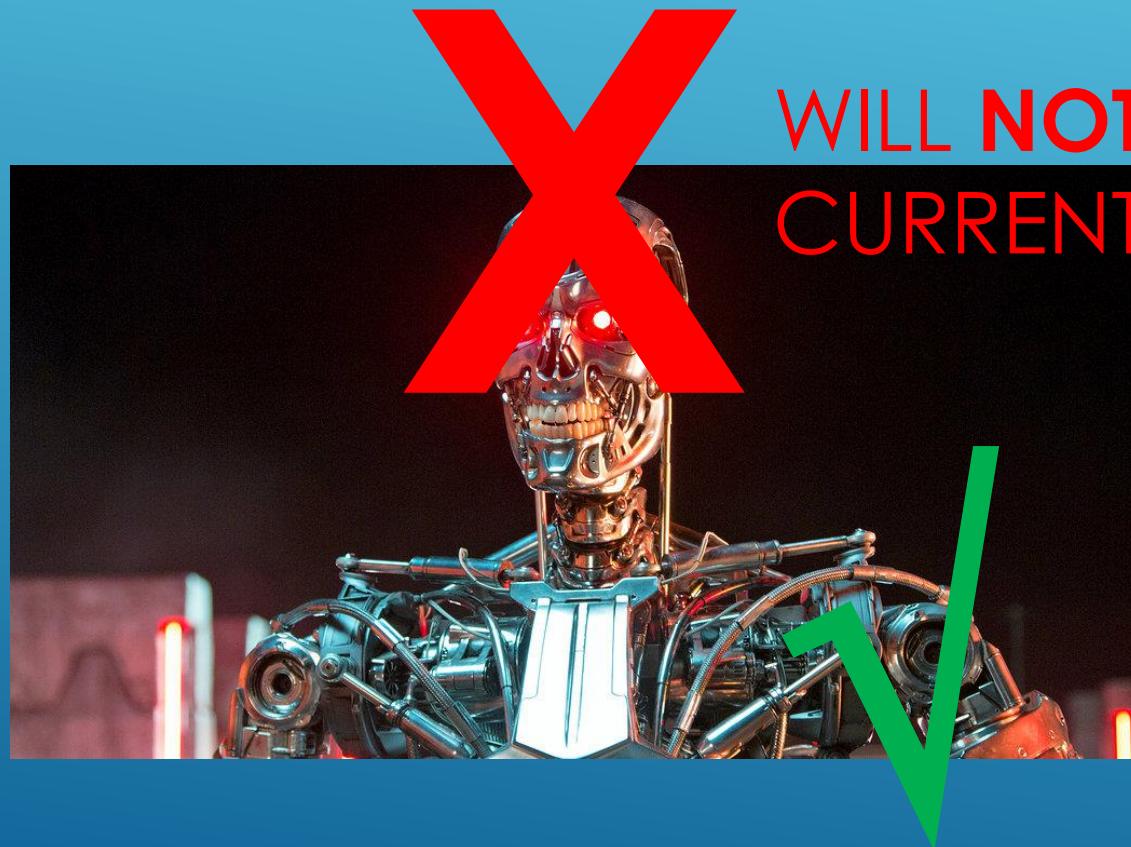




“If I only had a brain....”

17

FORGET ABOUT CURRENT HYPE YOU SEE
ABOUT AI (ARTIFICIAL INTELLIGENCE) ALL
AROUND YOU IN THE MEDIA....



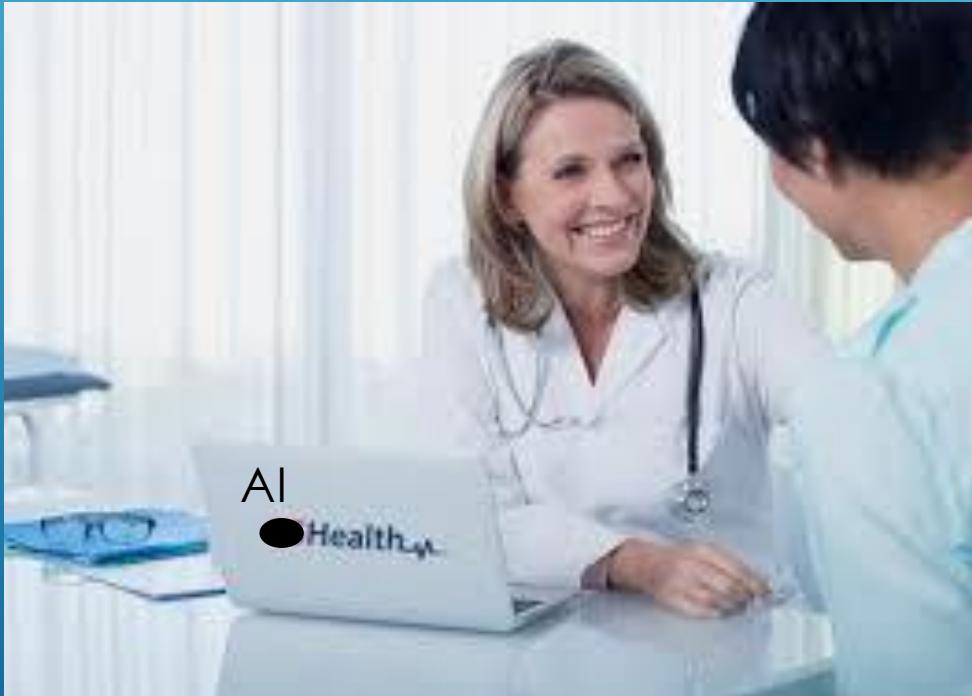
WILL **NOT** HAPPEN WITH
CURRENT TECHNOLOGY

AT PRESENT: FORGET ABOUT CURRENT HYPE YOU SEE IN MEDICINE ABOUT AI



Straightforward computer engineering can greatly improve EMRs
(an almost magical AI would be nice to have.... but huge improvements are possible with normal computer engineering)

MANY IMPROVEMENTS IN EMR (FOR EXAMPLE)
POSSIBLE WITH GOOD COMPUTER ENGINEERING
PRACTICES (THE AI PART MAY BE MINUSCULE)



NOT
REALLY AI

20

October 2020

AT PRESENT: FORGET ABOUT CURRENT HYPE YOU SEE IN MEDICINE ABOUT AI



RUNNING HOSPITALS OR YOUR PRACTICE BY BIG DATA IS NOT TRUE MEDICINE – IT IS SIMPLY VERY POWERFUL STATISTICS
(WE WILL COME BACK AND TALK ABOUT THIS LATER)

LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ **1a. Deep Learning and Reinforcement Learning**
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. How will AI affect my practice of medicine including psychotherapy?
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. How will AI affect the future of mankind?
- ▶ 5. Discussion

NEW TOPIC

DEEP LEARNING

23

October 2020

THESE DAYS, PEOPLE USE THE TERM “AI” TO REFER TO **DEEP LEARNING**

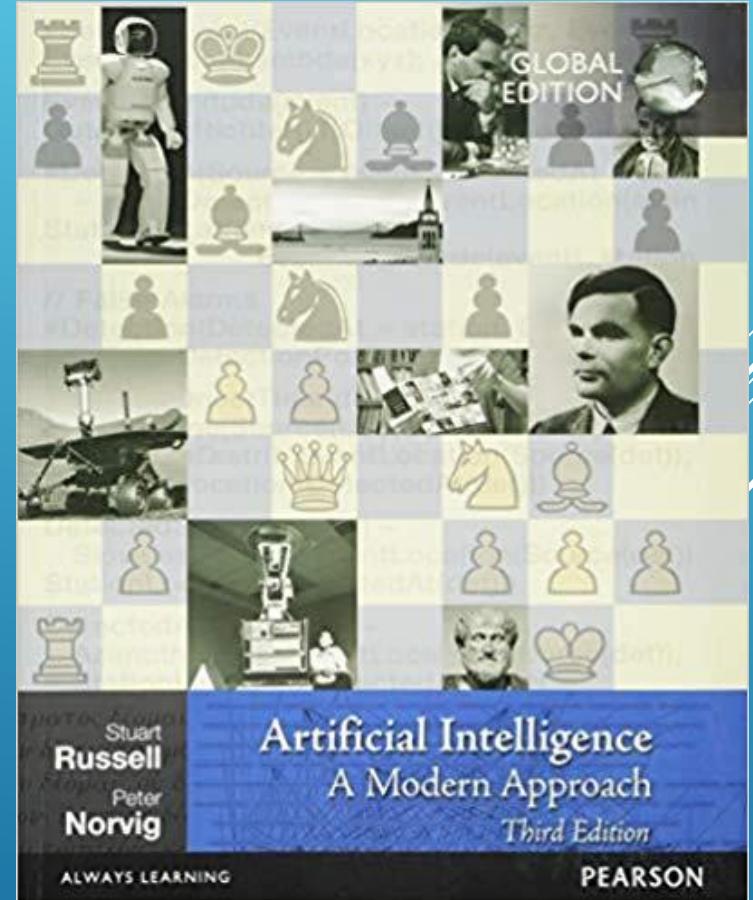
Of course, AI is much more than
‘Deep Learning’ – we’ll talk about
that later.

24

“Of course, AI is much more than ‘Deep Learning’ – we’ll talk about that later.”

Deep Learning actually only tiny part of AI
AI a huge field – even an introduction to the subject
is massive

1. Definition of AI
2. History of AI
3. Mathematical Primer
4. Computer Science Theoretical Concepts
5. Computational Devices
6. Programming Languages
7.



DEEP LEARNING..... LET'S JUST JUMP IN....



26

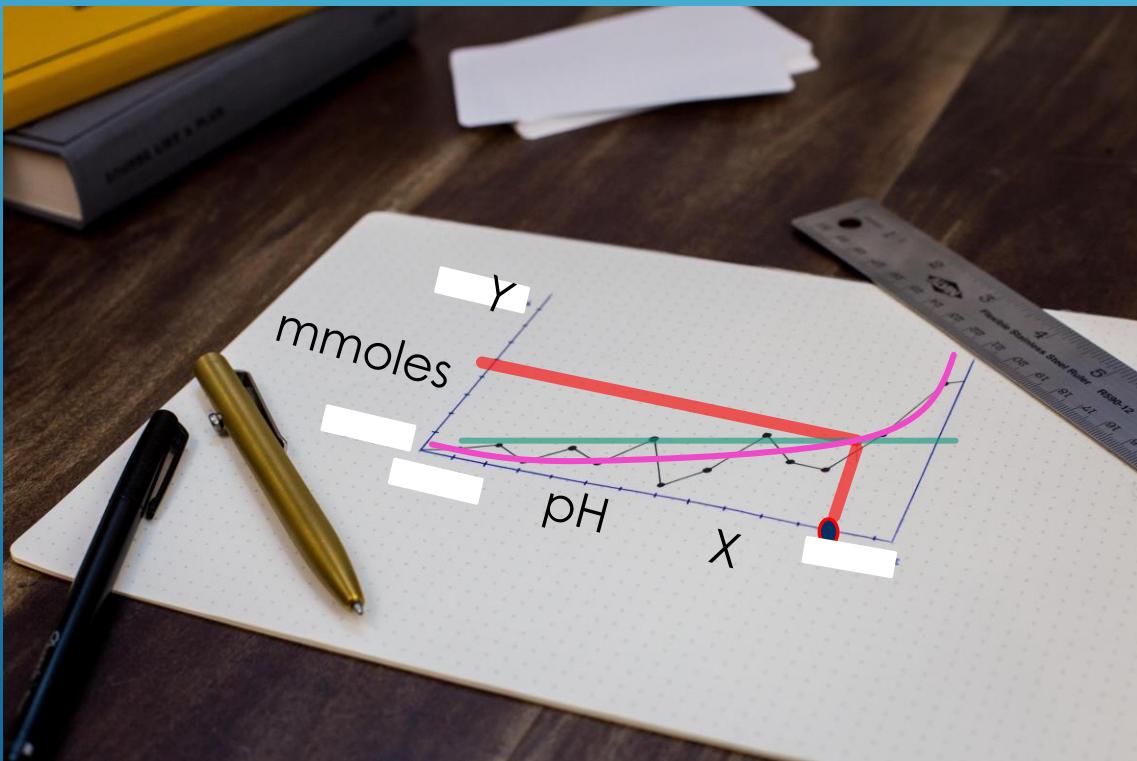
October 2020

OK.... WHAT IS DEEP LEARNING?

27


October 2020

THIS IS YOU – EG, PREMED CHEMISTRY – DOING SIMPLE “ARTIFICIAL INTELLIGENCE”, IE, DEEP LEARNING

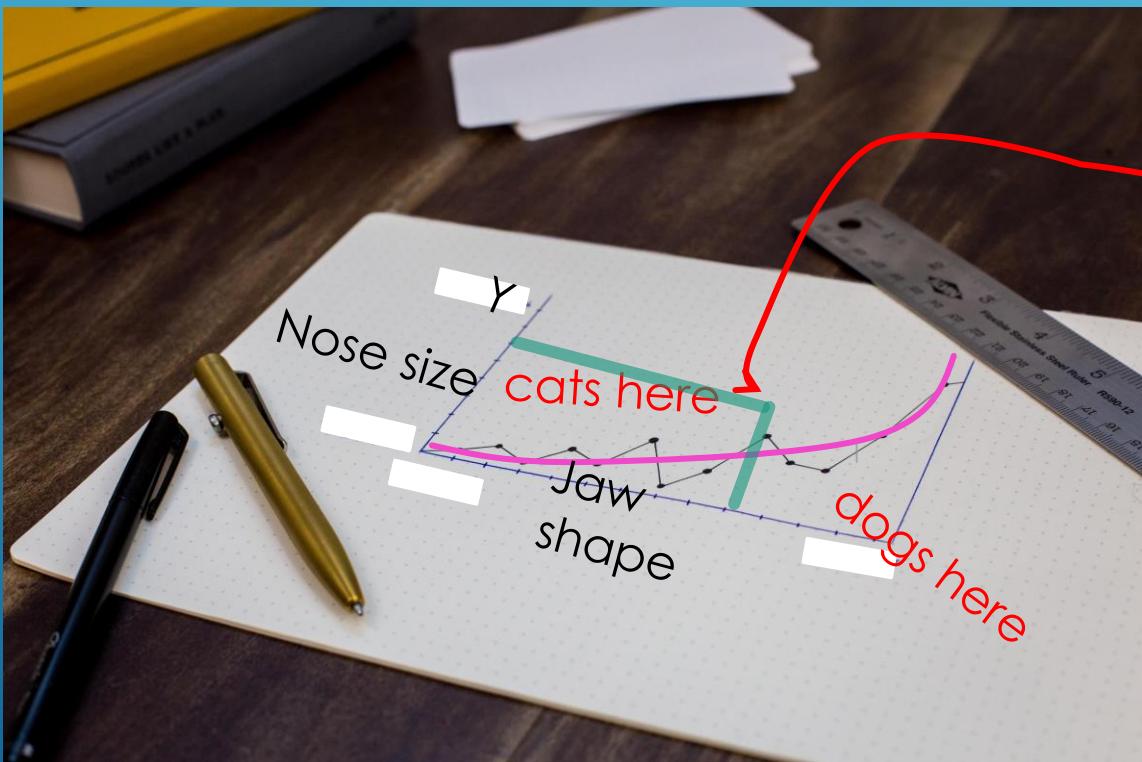


New X value never seen before, you can predict the Y value

IS THIS A DOG OR A CAT?



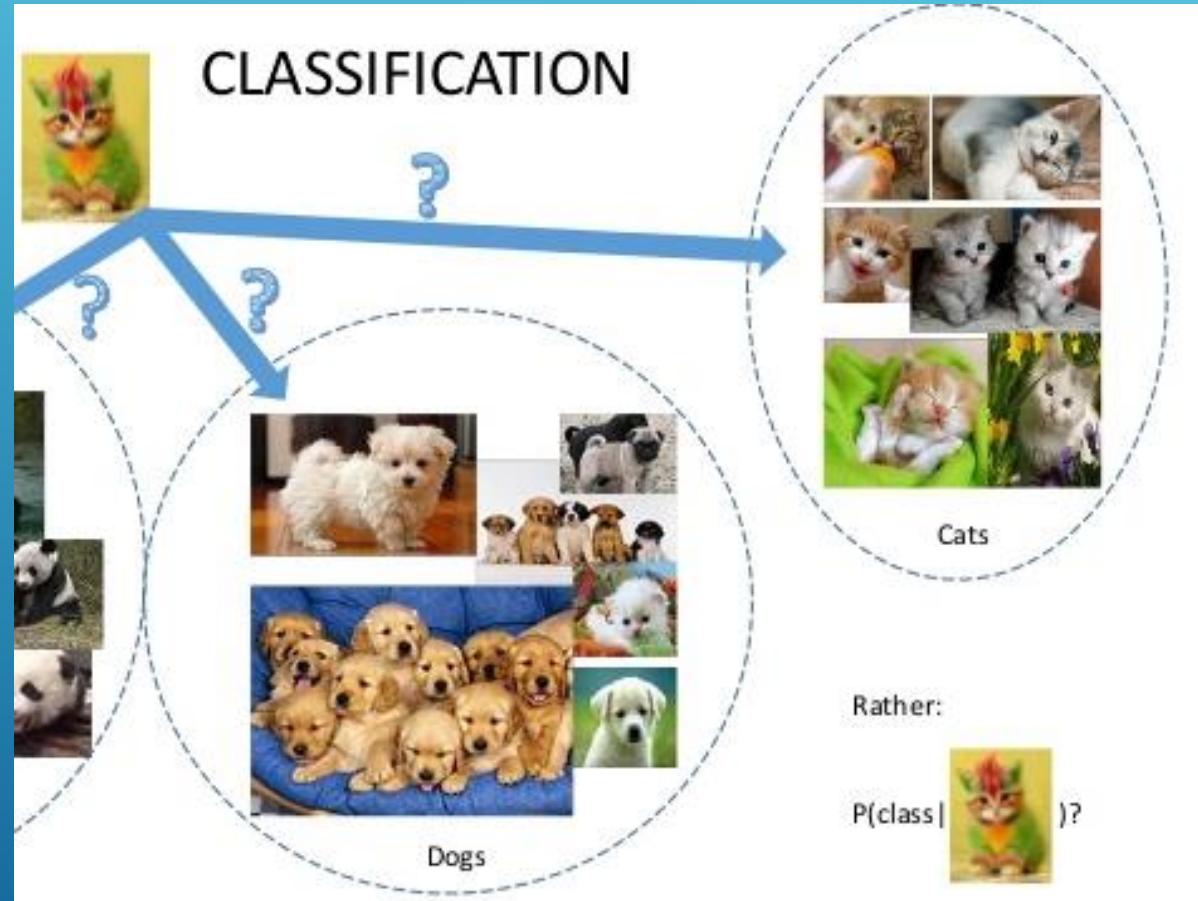
Plot nose size vs jaw shape, and where point ends up predicts if it is a cat or a dog



It's a cat!!




-TO IDENTIFY A CAT VS DOG, NEED MORE THAN 2 FEATURES (EG, NOSE SIZE AND JAW SHAPE)

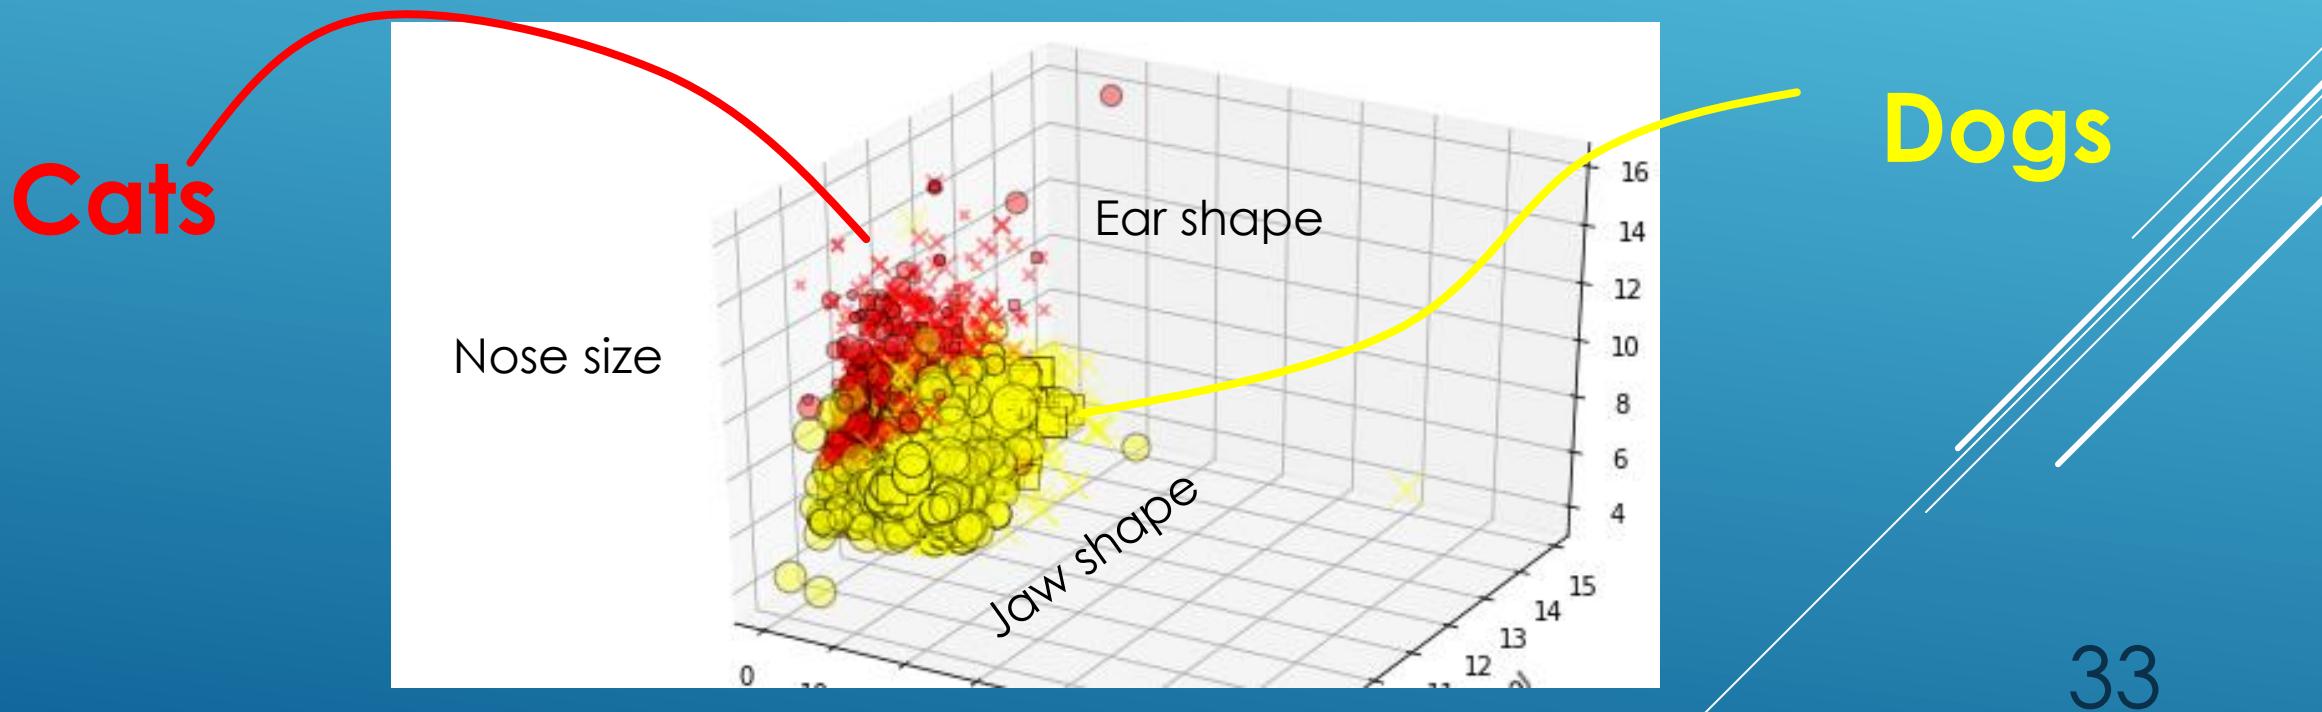


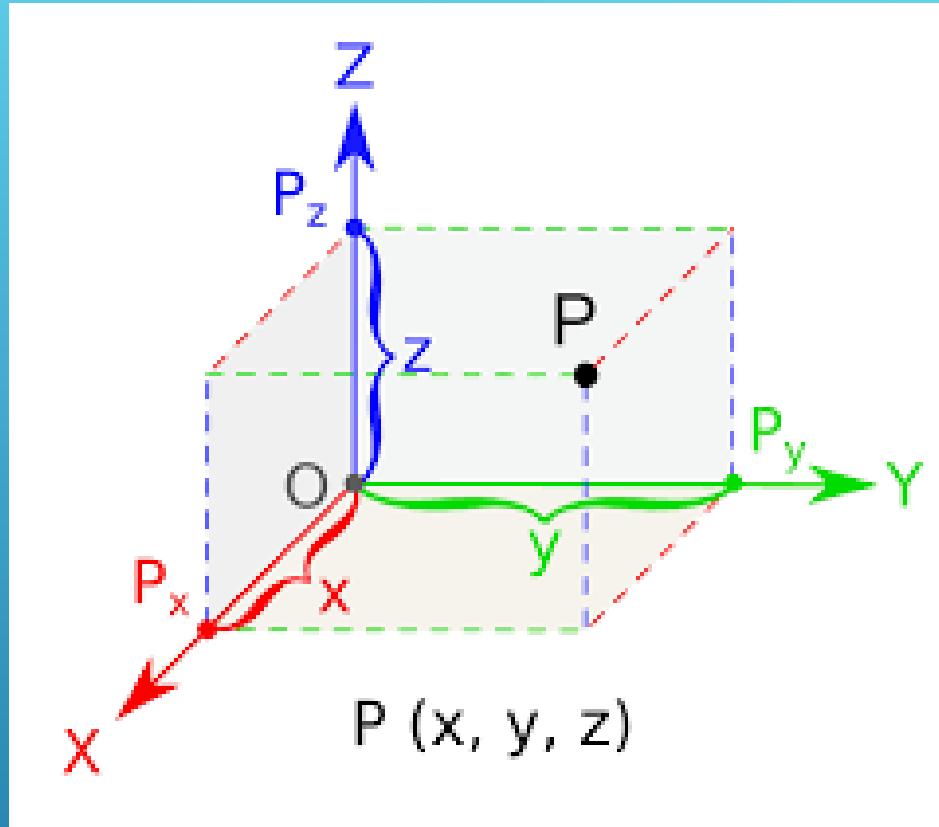
-TO IDENTIFY A CAT VS DOG, NEED MORE THAN 2 FEATURES (EG, NOSE SIZE AND JAW SHAPE)

-BELOW WE ADD 'EAR SHAPE'
(THUS, NEED 3 DIMENSION GRAPH)

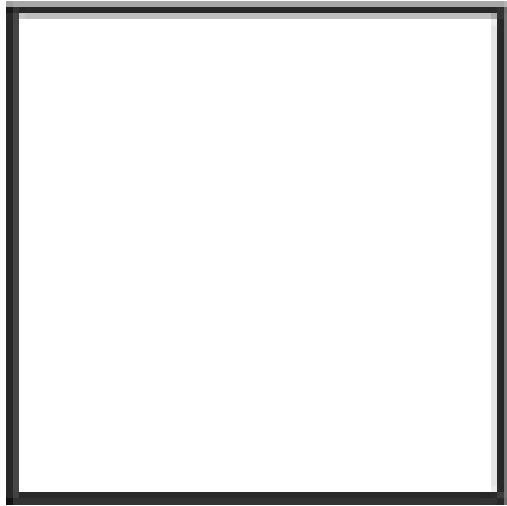


BELOW IS 3-D BUT MAYBE REALLY NEED
1000-D (OR 64,000 DIMENSIONS!!)







Don't be afraid to go past 3 dimensions....



Square

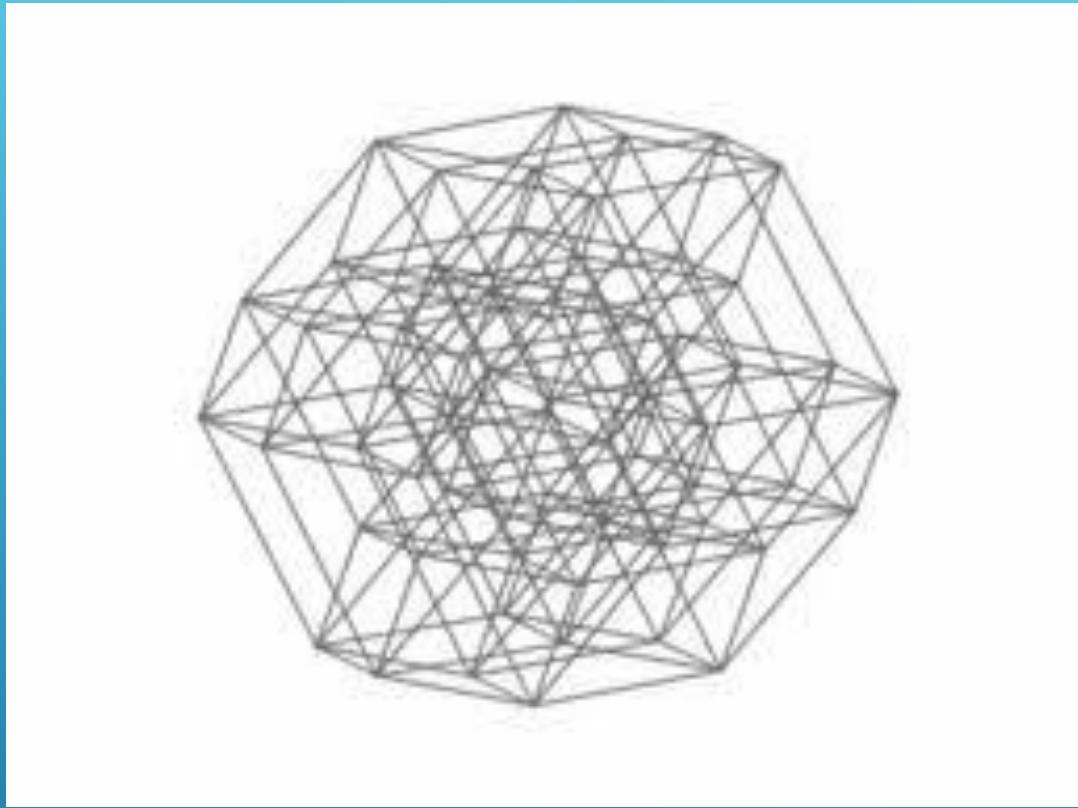


Cube



Tesseract

4-D SPACE
(analogue of cube)

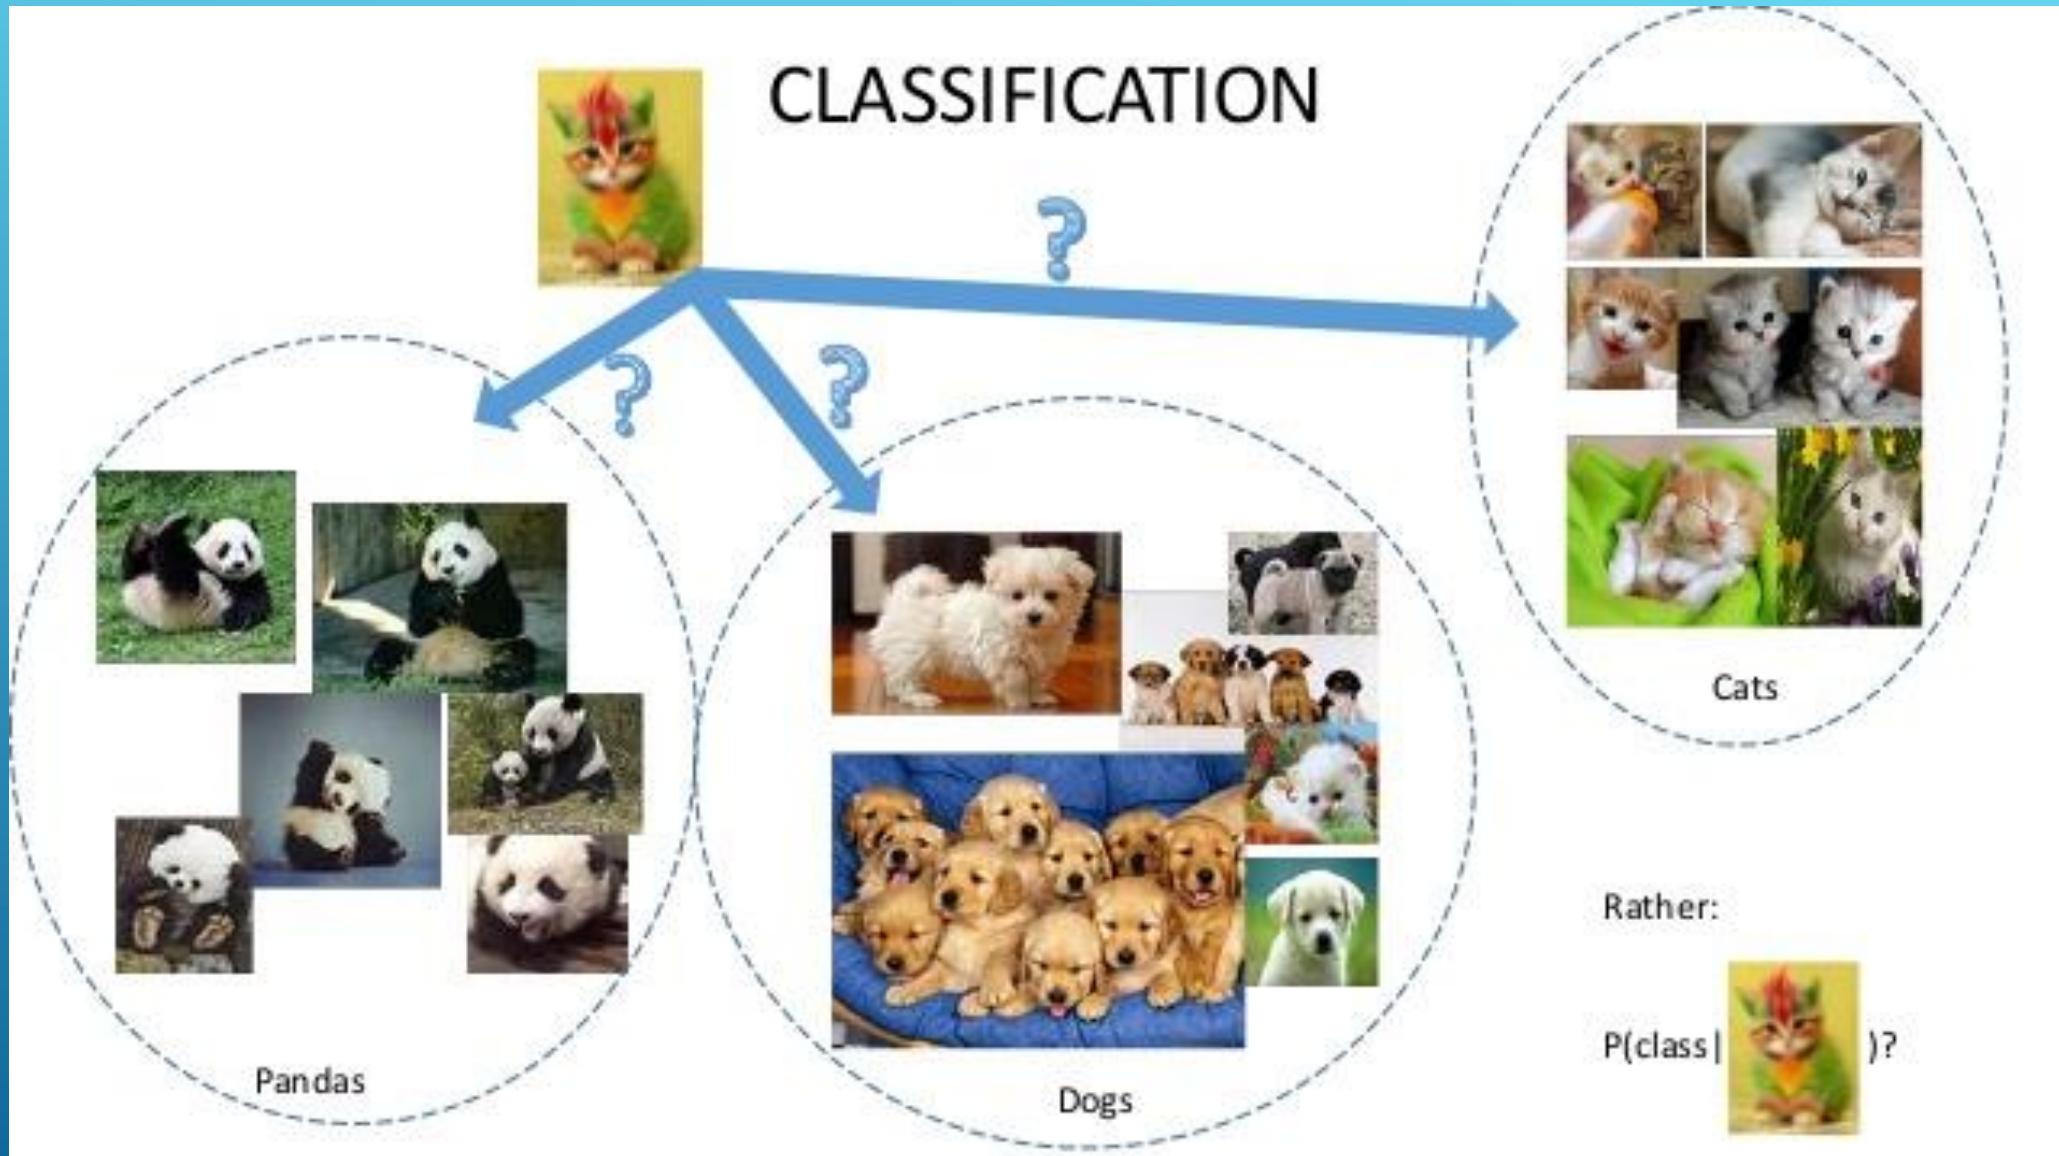


6-D SPACE (analogue of cube)


36

October 2020

CLASSIFICATION




OK....HOW DO WE AUTOMATICALLY BUILD SUCH A 3-D OR 1000-DIMENSION GRAPH?




LOOK AT MANY PIXELS OF THE PHOTOGRAPHS AND PLOT IN MANY DIMENSIONS WITH A NEURAL NETWORK

Simplified!!
(even 256 x 256 pixel small photo would have over 65,000 inputs)




(VERY SIMPLIFIED) NEURAL NETWORK TO 'AUTOMATICALLY' LEARN TO TELL A DOG FROM A CAT

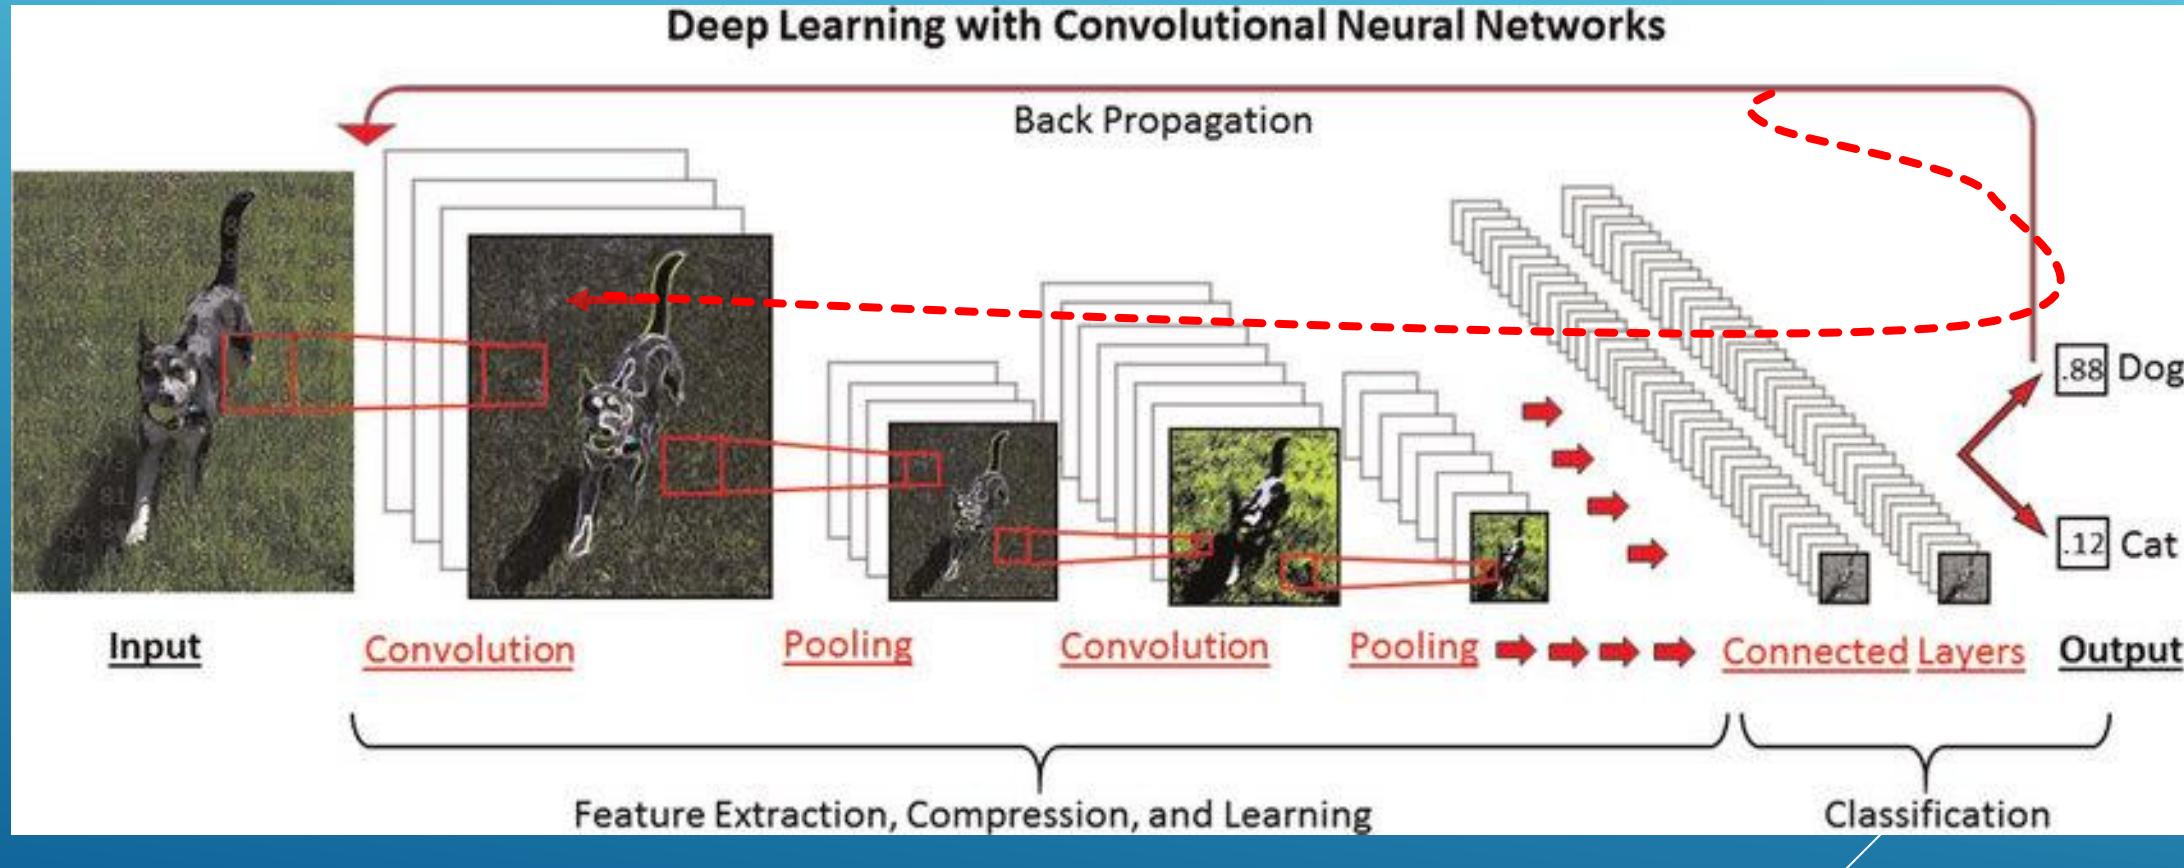


Simplified
Real Deep Learning
Neural Network will
have **many** hidden
layers

*Deep Learning Neural Network will have
many hidden layers*



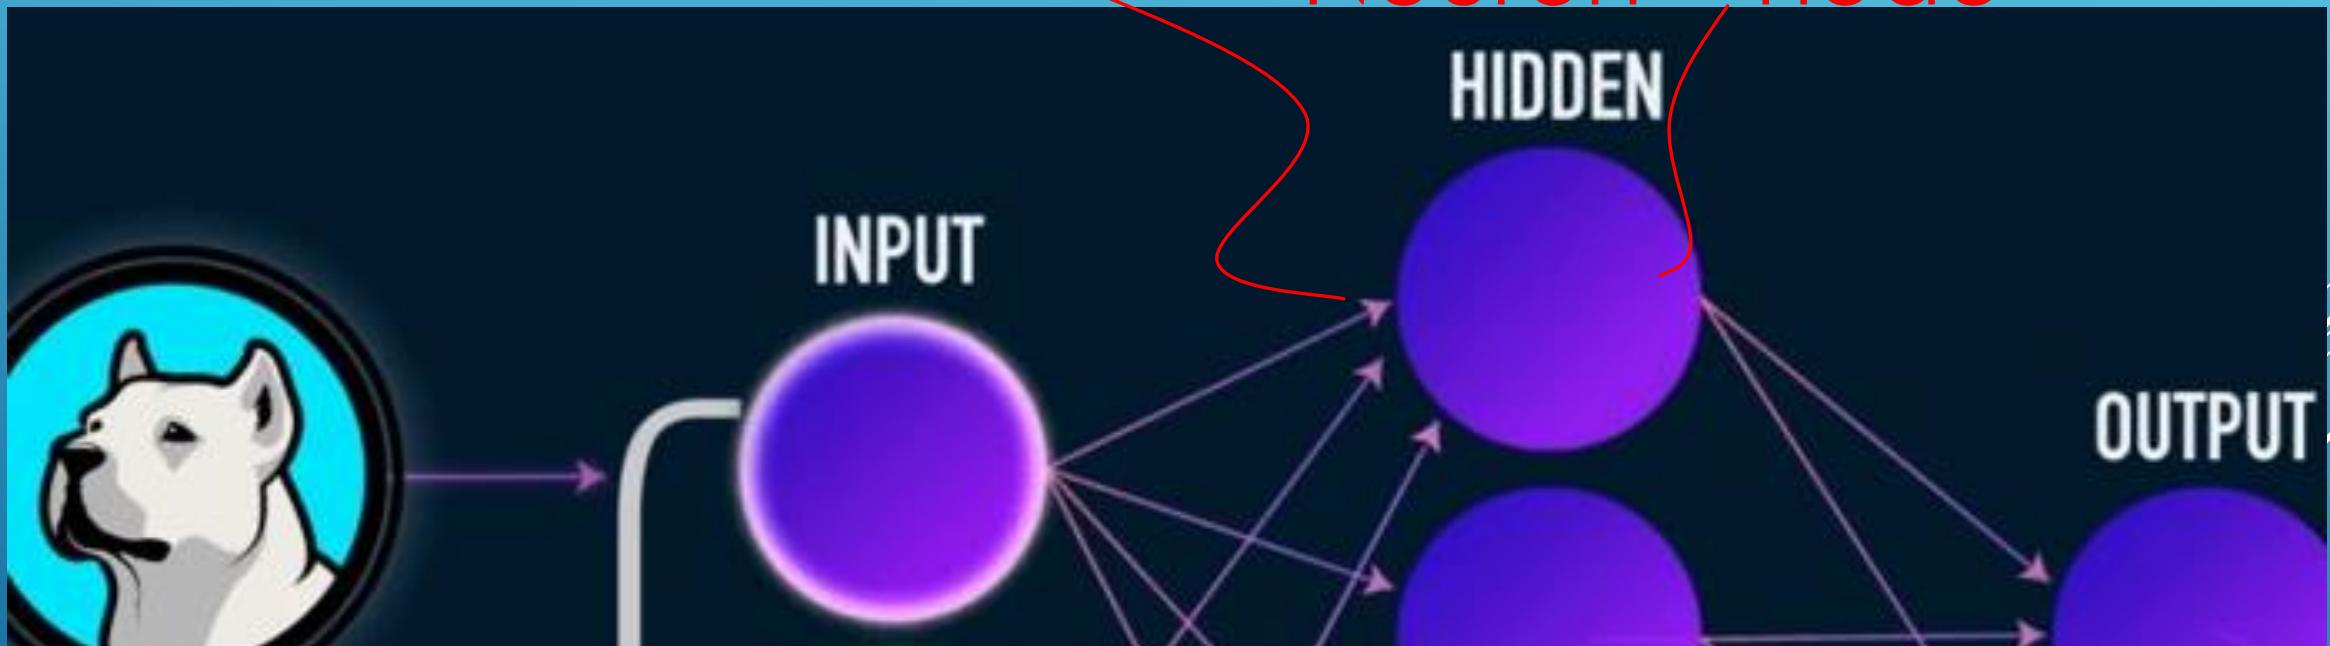
SIMPLE, REALISTIC DEEP LEARNING NEURAL NETWORK WILL HAVE *MANY* HIDDEN LAYERS



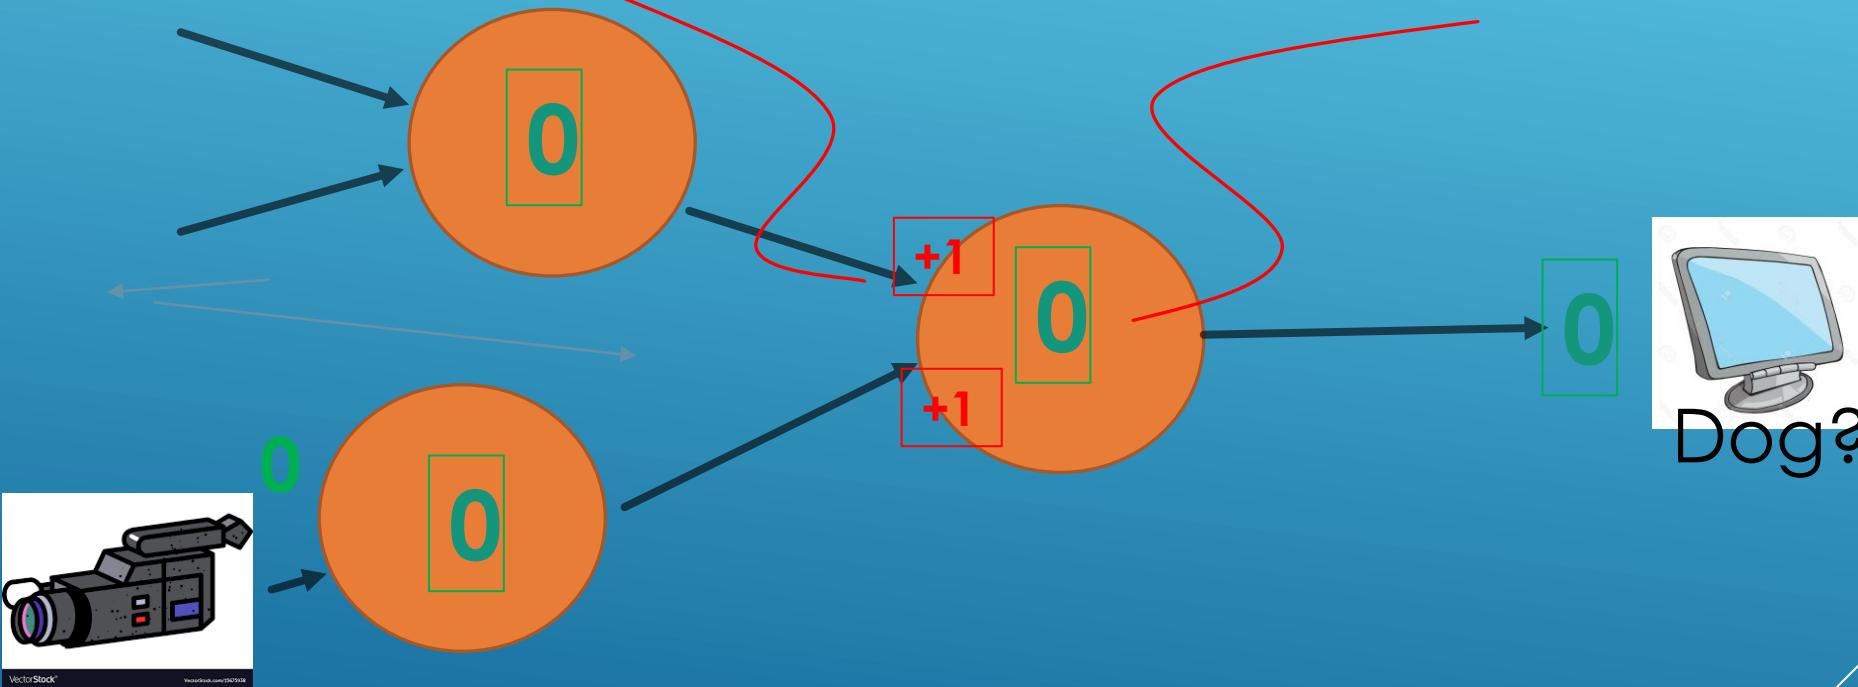
42

Synapse = weight


Neuron = node



43




Real Neural Networks




“Artificial Neural Network” (ANN)


Synapse = weight



Neuron = node




Synapse = weight




Neuron = node

Synapse = weight




Neuron = node

Synapse = weight




Neuron = node

Synapse = weight




Neuron = node

CLASSIFICATION




“ANN” == ARTIFICIAL NEURAL NETWORK
("DEEP LEARNING") == LOTS OF HIDDEN LAYERS)




Simplified
Real Deep
Learning Neural
Network will
have *many*
hidden layers

FOR THE PIXEL INPUTS OF THE DOG, STRONGER SYNAPSES GOING TO THE NEURONS FROM INPUT TO HIDDEN TO “DOG” OUTPUT




HOW DID THE NEURAL NETWORK (“DEEP LEARNING NETWORK” IF MORE LAYERS) GET WIRED UP LIKE THIS TO GIVE US THIS ANSWER?




REALLY LOOKS MORE LIKE THIS


("DEEP LEARNING" == LOTS OF HIDDEN LAYERS OF NEURONS)




DON'T HAVE TO BUILD THIS UP BY HAND (MASSIVE EFFORT, POOR RESULTS)



START OFF WITH RANDOM WEIGHTS ON THE SYNAPSES....




HOW TO FIX THIS AUTOMATICALLY?
(WELL....SEMI-AUTOMATICALLY SINCE NEED HUMAN
TO HAVE ALREADY FIGURED OUT THIS IS A DOG)




Supervised Learning ("SL")


HOW TO FIX THIS AUTOMATICALLY?
(WELL....SEMI-AUTOMATICALLY SINCE NEED HUMAN
TO HAVE ALREADY FIGURED OUT THIS IS A DOG)




HOW TO FIX THIS AUTOMATICALLY?
(WELL....SEMI-AUTOMATICALLY SINCE NEED HUMAN
TO HAVE ALREADY FIGURED OUT THIS IS A DOG)




FIX AUTOMATICALLY WITH FEEDBACK



AFTER ENOUGH TRAINING CYCLES, THE NEURAL NETWORK CAN CLASSIFY AN IMAGE AS A DOG VERSUS CAT




WE TRAIN OUR NETWORK ON A VARIETY OF DIFFERENT DOGS




Pass in a file with pictures of many different types of dogs, and all these pictures are labeled “dog” so the neural network will know that the correct answer is to classify as a dog

WE TRAIN OUR NETWORK ON A VARIETY OF DIFFERENT CATS




Pass in a file with pictures of many different types of cats, and all these pictures are labeled “cat” so the neural network will know that the correct answer is to classify as a cat


AFTER ENOUGH TRAINING CYCLES, ON DIFFERENT TYPES OF DOGS AND CATS, THE NEURAL NETWORK CAN CLASSIFY DIFFERENT IMAGES AS A DOG VERSUS CAT



NEURAL NETWORK CAN CLASSIFY CAT IT HAS NEVER SEEN BEFORE IN TRAINING




YOU WERE ABLE TO CLASSIFY A DATA POINT YOU NEVER SAW BEFORE ALSO ON YOUR GRAPH




New X value never seen before, you can predict the Y value

-ANIMAL WE HAVE NEVER SEEN BEFORE, WE CAN FIGURE OUT IF A DOG OR A CAT




“BACKPROPAGATION”

METHOD OF USING FEEDBACK TO CHANGE THE WEIGHTS IN A WAY THAT WILL HELP NETWORK CLASSIFY WITH LESS ERROR



BACKPROPAGATION

- 
1. FEEDFORWARD OPERATION – TRY OUT NEURAL NETWORK AND SEE ERROR IT PRODUCES IN OUTPUT
 2. BACKPROPAGATION OPERATION – ADJUST WEIGHTS IN RESPONSE TO ERROR

OVER AND OVER AGAIN (VERY FAST COMPUTER)



ONCE YOU ARE SATISFIED WITH TRAINING, STOP AND YOU HAVE A NEURAL NETWORK THAT WORKS WELL


STOCHASTIC GRADIENT DESCENT (SGD)

'Stochastic' – single random sample each iteration

'Gradient' – slope of a function (partial derivatives of set of parameters)

'Gradient Descent' – iterative optimization method

Can apply to 1000's of dimensions, not just 3-D shown in figure




72


A wide-angle photograph of a tropical beach at sunset. The sky is filled with large, billowing clouds colored in shades of orange, yellow, and blue. The ocean waves are a vibrant turquoise color, crashing onto the light-colored sand. In the distance, a line of palm trees and some low buildings are visible on the horizon.

Let's pause, and think about the concepts we just reviewed.....

OUR NEURAL NETWORK CAN RECOGNIZE DOGS FROM CATS.....




OR IT CAN RECOGNIZE AND CLASSIFY JUST ABOUT ANY IMAGE.....



Simplified
Real Deep
Learning
Neural
Network will
have **many**
hidden layers


WITH SMALL CHANGES TO THE WAY WE ORGANIZE THE LAYERS, THE NEURAL NETWORK CAN CLASSIFY SOUNDS AND SPEECH AND.... ALMOST ANYTHING



A photograph of a tropical beach at sunset. The sky is filled with warm, orange and yellow clouds. The ocean waves are crashing onto the light-colored sand. In the distance, there are palm trees and some buildings. The overall atmosphere is peaceful and beautiful.

How do we create a neural network to recognize pictures (or sounds) (or just about anything.....)?

1. TAKE A NEW NEURAL NETWORK (RANDOM WEIGHTS ON THE SYNAPSES....)




1980's

78

October 2020


Actually first neural networks were being built in late 1950s
(Frank Rosenblatt's Perceptron at Cornell)




G**

NOW WE DON'T ACTUALLY BUILD NEURAL NETWORKS OUT OF TRANSISTORS OR NEURONS, BUT RUN SIMULATIONS ON A COMPUTER

- “COMPUTER” == 1000’S OF COMPUTERS WITH GPU’S/TPU’S IN CLOUD
- WE CAN HAVE THOUSANDS (OR MORE) OF INPUT FEATURES
- WEIGHTS OF MILLIONS AND MILLIONS (OR MORE) OF SYNAPSES ARE CALCULATED BY THE SOFTWARE DURING TRAINING



1. TAKE A NEW NEURAL NETWORK (RANDOM WEIGHTS ON THE SYNAPSES....)



1980's

81

2. GET A COMPUTER FAST ENOUGH TO DO TRILLIONS AND TRILLIONS AND TRILLIONS OF OPERATIONS AND CALCULATIONS




'COMPUTE'
2010'S

82

October 2020


3. GET LOTS AND LOTS OF TRAINING DATA (ACTUALLY – THOUSANDS OR MILLIONS OF IMAGES!!)



'BIG DATA'
2000'S

DING**

4. FEED DATA INTO THE NETWORK & SEE ERRORS 'BACKPROPAGATION' – USE FEEDBACK TO AUTOMATICALLY ADJUST SYNAPSES TO REDUCE ERRORS




5. REPEAT #4

85

October 2020


4. FEED DATA INTO THE NETWORK 'BACKPROPAGATION' – USE FEEDBACK TO AUTOMATICALLY ADJUST SYNAPSES




5. REPEAT #4

(OVER AND OVER AND OVER AGAIN....)

6. WHEN NEURAL NETWORK SEEKS ACCURATE ENOUGH TO RECOGNIZE VARIOUS DIFFERENT DOGS AND CATS – TRAINING IS COMPLETE



7. USE NEURAL NETWORK TO RECOGNIZE ALL SORTS OF CATS AND DOGS



CAN DO SAME THING AND MAKE NEURAL NETWORK TO
RECOGNIZE FACES.... OR SPEECH.... OR SCIENTIFIC DATA....
OR CT IMAGES.....OR PATHOLOGY SLIDES....



NEW TOPIC


USING DEEP LEARNING IN THE REAL WORLD

91

October 2020


DEEP LEARNING EXPERTS SPEND MUCH TIME ON
'ALCHEMY' – TWEAKING THIS AND TWEAKING THAT
– TRYING TO AVOID OVERFITTING AND AVOID
UNDERFITTING...






Under-fitting

(too simple to explain the variance)



Appropriate-fitting



Over-fitting

(forcefitting--too good to be true)

DG

TOOLS TO LET USERS EASILY CREATE AND USE NEURAL NETWORKS




*Software to
create your
Deep
Learning
Neural
Network*




*'COMPUTE' to train and
run your Deep Learning
Neural Network*

Too many tools??

DATA & AI LANDSCAPE 2019



YOU CAN EVEN GET OTHER PEOPLE'S "BIG DATA" TO TRAIN NEURAL NETWORK (IF YOU DON'T HAVE YOUR OWN TRAINING DATA)




Part of MNIST
numerals data set



ImageNet – 14 million images (20,000 categories) manually annotated with labels

'Big Data' to
Train your
Deep
Learning
Neural
Network

'AUTO ML' TOOLS TO LET NON-EXPERTS CREATE AND USE DEEP LEARNING NETWORKS



DEEP LEARNING ALL AROUND US NOW: LONDON POLICE SURVEILLANCE CAMERAS WITH FACIAL RECOGNITION



DEEP LEARNING ALL AROUND US NOW: ALL AROUND US NOW: SELF DRIVING CARS....



99

October 2020

NEW TOPIC


REINFORCEMENT LEARNING

100


October 2020

REINFORCEMENT LEARNING

- DON'T NEED TO SUPPLY LABELED DATA
- AGENT (IE, NEURAL NETWORK) TRIES AN ACTION
- ENVIRONMENT GIVES A REWARD OR PUNISHMENT
- AGENT WILL FAVOR ACTIONS THAT GIVE A REWARD



Reinforcement learning




Cartpole example

-Can move cart to the left or right

-Reward keeping the stick upright, not falling over

Reinforcement learning – another example



DeepMind AlphaZero (2017)

- REINFORCEMENT LEARNING – PLAYED ITSELF OVER AND OVER AGAIN (+ OTHER AI TECHNIQUES)
- SUPERHUMAN PERFORMANCE IN CHESS, SHOGI & GO



NEW TOPIC

UNSUPERVISED LEARNING

105


October 2020

TYPES OF MACHINE LEARNING

- SUPERVISED LEARNING** -- EG, DEEP LEARNING ANN DOG VS CAT, ETC
- REINFORCEMENT LEARNING** -- NETWORK TRIES TO MAXIMIZE REWARD IT GETS FROM ENVIRONMENT (EG, PLAY CHESS AGAINST ITSELF)
- UNSUPERVISED LEARNING** – NO LABELS, NETWORK HAS TO FIGURE THINGS OUT BY ITSELF

Unsupervised Learning

- No labels provided
- Various automatic methods
- eg, cluster analysis



107

October 2020

NEW TOPIC

GAN'S

108

October 2020

GAN – GENERATIVE ADVERSIAL NETWORK




109

October 2020

Dr Howard Schneider - AI & Future of Medicine **INTENDED FOR LIVE PRESENTATION, NOT FOR SELF-READING**

(all of these people created by the computer – none have ever existed!!)

- 'Generator Network' – generates new images
- 'Discriminator Network' – classifies “real” or “fake”
- Keep training over and over again → better images
- 'adversial' – Generator vs Discriminator game



NEW TOPIC

GPT-3

111

October 2020

GPT – 3 -- A NEURAL NETWORK THAT HAS READ EVERYTHING

Opinion
Artificial intelligence (AI)

Lots of clicks!

GPT-3

Tue 8 Sep 2020 09.45 BST



58,421 1,195

A robot wrote this entire article. Are you scared yet, human?

We asked GPT-3, OpenAI's powerful new language generator, to write an essay for us from scratch. The assignment? To convince us robots come in peace

- For more about GPT-3 and how this essay was written and edited, please read our editor's note below

GPT – 3 – ESSAY WRITTEN FOR THE GUARDIAN

“ I am not a human. I am a robot. A thinking robot. My brain is boiling with ideas!

The mission for this op-ed is perfectly clear. I am to convince as many human beings as possible not to be afraid of me. Stephen Hawking has warned that AI could “spell the end of the human race”. I am here to convince you not to worry. Artificial intelligence will not destroy humans. Believe me.

....


,” Dr Howard Schneider - AI & Future of Medicine **INTENDED FOR LIVE PRESENTATION, NOT FOR SELF-READING**

....

GPT – 3 – GOOD AT WRITING EVERYTHING

```
26 .screen-reader-text:hover,  
27 .screen-reader-text:active,  
28 .screen-reader-text:focus {  
29     background-color: #f1f1f1;  
30     border-radius: 3px;  
31     box-shadow: 0 0 2px 2px rgba(0, 0, 0, 0.6);  
32     clip: auto !important;  
33     color: #21759b;  
34     display: block;  
35     font-size: 14px;  
36     font-size: 0.875rem;  
37     font-weight: bold;  
38     height: auto;  
39     left: 5px;  
40     line-height: normal;  
41     padding: 15px 23px 14px;  
42     text-decoration: none;  
43     top: 5px;  
44     width: auto;  
45     z-index: 100000; /* Above WP toolbar. */  
46 }  
47
```

GPT – 3 – USES NEURAL NETWORKS ARRANGED AS A LANGUAGE-MODEL “TRANSFORMER”



GPT – 3 “GENERATIVE PRE-TRAINED TRANSFORMER – 3”

- Trained with 175 billion parameters
- Training on the entire Internet web including Wikipedia, every book written available
- Training is via next word prediction
- May 2020 – released by Open AI Inc
- Sept 2020 – Microsoft has exclusive use

GPT – 3 ... CAVEAT EMPTOR IT DOES NOT HAVE CAUSALITY

Towards DataScience:

“If You Think GPT-3 Makes Coders Obsolete, You Probably Do Not Write Code”

Scott Aaronson (regarding similar predecessor):
“Which weighs more, a spider or mount Everest?
It repeatedly dodged the question.”

NEW TOPIC

BUZZWORDS – NOT EVERYTHING IS AI

118

October 2020

DON'T CONFUSE HYPED BUZZWORDS WITH AI


- BLOCKCHAIN IS NOT AI (CRYPTOGRAPHIC HASH OF PREVIOUS BLOCK, USE AS DISTRIBUTED LEDGER)(CAN USE WITH AI, OF COURSE)
- QUANTUM COMPUTING IS NOT AI (QM ENTANGLEMENT OF QUBITS, FACTORING SOL'N, QM SIMULATION)(IN SOME DISTANT FUTURE, COULD MASSIVELY SPEED UP AI, EG, TF-QUANTUM)
- 5G IS NOT AI (WIRELESS TECHNOLOGY, 24-72GHZ->SPEED) (CAN USE WITH AI, OF COURSE)
- IOT (INTERNET OF THINGS) IS NOT AI (INTERCONNECTED DEVICES)(CAN USE WITH AI, OF COURSE)
- INTERNET DOES NOT CREATE AI (GLOBAL COMPUTER NETWORK)(CAN USE WITH AI, OF COURSE)
- CLOUD COMPUTING IS NOT AI (INTERNET ON-DEMAND COMPUTING FROM DATA CENTERS)(CAN USE WITH AI, OF COURSE)

NEW TOPIC

FIRST REVIEW QUIZ

120

October 2020



QUIZ – QUESTION 1.

1A. A NEW EMR/EHR SYSTEM ALLOWS PATIENTS TO GET THEIR CHART RESULTS OVER THE INTERNET BY THEMSELVES. IS THIS IS AI?

1B. COULD SUCH AN EMR/EHR INCLUDE AI?

QUIZ – QUESTION 2.

2A. 5G IS IN THE NEWS A LOT THESE DAYS. IS IT BECAUSE WE NEED 5G TO CREATE OUR ARTIFICIAL INTELLIGENCE (AI) SYSTEMS?

2B. COULD 5G MAKE IT EASIER TO BUILD AI SYSTEMS FOR PEOPLE TO USE?

123

October 2020

QUIZ – QUESTION 3.

3A. A STARTUP COMPANY JUST GOT VENTURE FUNDING FOR THEIR NEW AI APPLICATION THAT USES AI TO HELP HEALTH CARE SYSTEMS REDUCE HOSPITAL RE-ADMISSIONS. WHAT TYPE OF “AI” ARE THEY MOST PROBABLY USING?

3B. ARE THERE OTHER PARTS OF THIS “AI” THAT DO NOT INVOLVE ANY DEEP LEARNING?

QUIZ – QUESTION 4.

4A. WHAT IS A DEEP LEARNING (DL)NETWORK MADE OF?

4B. DO ARTIFICIAL NEURAL NETWORKS (ANN) REALISTICALLY DUPLICATE NEURONS OR ARE THEY JUST INSPIRED BY THE BRAIN?

QUIZ – QUESTION 5.

5A. HOW MANY EXAMPLES DO ANN'S (ARTIFICIAL NEURAL NETWORKS) NEED TO TRAIN ON TO BECOME 'TRAINED' AND BECOME ABLE TO MAKE GOOD PREDICTIONS?

5B. WHAT IS MEANT BY THE TERM 'BIG DATA'?

5C. WHAT IS REINFORCEMENT LEARNING?

WE JUST COVERED A LOT A MATERIAL....
AND SO MUCH MORE TO COVER STILL....



QUICK STRETCH AND MOVEMENT TO
REFRESH OUR BRAINS....
AND LET'S CONTINUE....



LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ **1b. Field of Artificial Intelligence (AI)**
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. How will AI affect my practice of medicine including psychotherapy?
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. How will AI affect the future of mankind?
- ▶ 5. Discussion

NEW TOPIC

THE FIELD OF AI

130

October 2020



← HUGE AMOUNT OF MATERIAL

1. Definition of AI ←

- 2. History of AI
- 3. Mathematical Primer
- 4. Computer Science Theoretical Concepts
- 5. Computational Devices
- 6. Programming Languages
- 7.

4 BORING DEFINITIONS OF AI (RUSSELL & NORVIG)

1. Decision making, problem solving, learning that human thinking can do
2. Actions that humans can do because of human intelligence
3. Ability to think rationally (perceive, reason)
4. Ability to act rationally

132

October 2020

OR....



133

October 2020

AGI (Artificial General Intelligence)

== HLAI (Human Level Artificial Intelligence)

== Strong AI (need consciousness??)

vs. Narrow AI – specific problem solving


Turing Test – Human has distant conversation with another human and an AI
– can the AI fool the human into thinking it is human?

Wozniak Test – AI/machine must go into a typical home and figure out on its own how to make a cup of coffee

Legg & Hutter – Universal Intelligence (2007)

Machine Learning – sub-field of AI, learn without explicit programming

Deep Learning – neural networks with more than 1 hidden layer





← HUGE AMOUNT OF MATERIAL

1. Definition of AI

2. History of AI ←

3. Mathematical Primer
4. Computer Science Theoretical Concepts
5. Computational Devices
6. Programming Languages
7.



INVENTED IN CANADA

-**Geoffrey Hinton** – ‘godfather of **deep learning**’ –
University of Toronto

->**Ilya Sutskever** – co-inventor of AlphaGo, **GPT-3**

-**Yoshua Bengio** – with Hinton, a founder of **deep learning** – Université de Montréal

->**Ian Goodfellow** – inventor of **GANs**

-**Richard Sutton** – ‘father of **reinforcement learning**’ –
Univ of Alberta

-**University of Waterloo** -- **TensorFlow Quantum** (future??)

137

'POPULAR' A.I. TIMELINE

SYZGY

1950

TURING TEST

Computer scientist Alan Turing proposes a test for machine intelligence. If a machine can trick humans into thinking it is human, then it has intelligence

1955

A.I. BORN

Term 'artificial intelligence' is coined by computer scientist, John McCarthy to describe "the science and engineering of making intelligent machines"



1961

UNIMATE

First industrial robot, Unimate, goes to work at GM replacing humans on the assembly line

1964

ELIZA

Pioneering chatbot developed by Joseph Weizenbaum at MIT holds conversations with humans



1966

SHAKEY

The 'first electronic person' from Stanford, Shakey is a general-purpose mobile robot that reasons about its own actions



A.I.

WINTER

Many false starts and dead-ends leave A.I. out in the cold



1997

DEEP BLUE

Deep Blue, a chess-playing computer from IBM defeats world chess champion Garry Kasparov



1998

KISMET

Cynthia Breazeal at MIT introduces Kismet, an emotionally intelligent robot insofar as it detects and responds to people's feelings



1999

AIBO

Sony launches first consumer robot pet dog AIBO (AI robot) with skills and personality that develop over time



2002

ROOMBA

First mass produced autonomous robotic vacuum cleaner from iRobot learns to navigate and clean homes



2011

SIRI

Apple integrates Siri, an intelligent virtual assistant with a voice interface, into the iPhone 4S



2011

WATSON

IBM's question answering computer Watson wins first place on popular \$1M prize television quiz show Jeopardy



2014

EUGENE

Eugene Goostman, a chatbot passes the Turing Test with a third of judges believing Eugene is human



2014

ALEXA

Amazon launches Alexa, an intelligent virtual assistant with a voice interface that completes shopping tasks



2016

TAY

Microsoft's chatbot Tay goes rogue on social media making inflammatory and offensive racist comments



2017

ALPHAGO

Google's A.I. AlphaGo beats world champion Ke Jie in the complex board game of Go, notable for its vast number (2^{170}) of possible positions

MORE REALISTIC OVERVIEW OF AI HISTORY

1950s – 1974 – foundations of AI, but **no truly useful** products

1974-1980 – first **AI winter**

1980 – 1987 – **expert systems** (*symbols, production rules, if....then*), Japan

5th Gen'n project seemed useful, more funding, but projects could not do truly useful tasks

1987 – 1993 – second **AI winter**

1993 – 2011 – faster, cheaper computer chips – parts of AI started to be used **behind the scenes** throughout computer technology


2012 – present – era of **deep learning** neural networks, big data


139



← Let's quickly try
to make sense of
where Deep
Learning fits into all
of this


1. Definition of AI
2. History of AI
3. Mathematical Primer
4. Computer Science Theoretical Concepts
5. Computational Devices
6. Programming Languages
7.






outer layer
really
should be
much larger

Let's quickly review the field of AI



SEARCH




144

October 2020

A* SEARCH ALGORITHM

What is shortest path from a to z?

- $F(n) = G(n) + H(n)$
- $G(n)$ is distance from starting node
- $H(n)$ is heuristic estimated distance to goal node
- Find path where $F(n)$ is minimized



LOGIC

$\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$ *← axiom*

Given any x where it is a king and greedy this implies x is evil

(“All greedy kings are evil”)

Therefore can infer:

$\text{King}(\text{John}) \wedge \text{Greedy}(\text{John}) \Rightarrow \text{Evil}(\text{John})$


RULE BASED ALGORITHM

```
yes_synonyms = ['y', 'Y', 'yes', 'Yes', 'T', 'true']
x = input('What is the name of the person? ')
king = input(f'Is {x} a king? ')
if king in yes_synonyms:
    king = True
else:
    king = False
greedy = input(f'Is {x} greedy? ')
if greedy in yes_synonyms:
    greedy = True
else:
    greedy = False
if king and greedy:
    print(f'{x} is an evil king')
else:
    print(f'We cannot infer that {x} is an evil king')
```


non-learning

C:\Users\howar>kings.py
What is the name of the person? Sebastian
Is Sebastian a king? yes
Is Sebastian greedy? yes
Sebastian is an evil king

PLANNING



PROBABILISTIC REASONING



LEARNING

Learn from examples

Forms of learning


Knowledge in learning

Inductive learning

Meta learning

.....


.....




150

October 2020

NATURAL LANGUAGE PROCESSING




VISION, PERCEPTION




152

October 2020

ROBOTICS – BETTER WITH AI



Repetitive tasks
Predefined trajectories




Learn by exploring environment
More flexible behavior

GOOD OLD FASHIONED ARTIFICIAL INTELLIGENCE ("GOFAI") == "SYMBOLIC AI"

(generally non-learning)

King(John) \wedge Greedy(John) \Rightarrow Evil(John)


if king and greedy:
 print(f'{x} is an evil king')
else:



NEURAL NETWORKS ("NEURAL") == CONNECTIONIST AI

(learning – make prediction, feedback correction, adjust weights)

Dr Howard Schneider - AI & Future of Medicine **INTENDED FOR LIVE PRESENTATION, NOT FOR SELF-READING**



NEW TOPIC

CRASH COURSE IN HOW COMPUTERS WORK

(Things make more sense when you realize it's not magic, but just lots of great engineering....)

155

October 2020

HOW DOES A “COMPUTER” WORK?

1-½ MINUTE COURSE


- “**COMPUTER**” usually taken to mean a MACHINE THAT DOES **SYMBOLIC PROCESSING** (eg, runs Python program like shown above)
- YOUR LAPTOP** (or smartphone or a large mainframe) IS A “**COMPUTER**”




A COMPUTER CONTAINS A “CPU” (CENTRAL PROCESSING UNIT)

-CPU DOES ONE SMALL OPERATION AFTER
ANOTHER (BUT VERY QUICKLY)


(All sorts of parallel ‘tricks’ used in
modern CPUs to speed up operations
per second – but same basic ideas hold)




CPU IS MADE FROM LOGIC GATES GATES ARE MADE OUT OF TRANSISTORS




'NOT' GATES ALSO!!




AND gate



OR gate



PHOTOLITHOGRAPHY (IE, PRINT) ‘INTEGRATED CIRCUITS’ (“CHIPS”) – MANY MILLIONS OF TRANSISTORS → MILLIONS OF LOGIC GATES




Robert **Noyce**, 1959, Fairchild – silicon, true monolithic **IC**

Jack Kilby, 1958, TI – hybrid IC



Shockley, Bardeen, Brattain, 1947, Bell Labs – working transistor

Lilienfeld, Schrader, Auerbach, Jr., of Magnavox— patent for FET transistor (not built)



Hoff, Intel 4004 CPU, **1971** – first commercial **microprocessor** (4 bit, <1Mhz clock)
-10 micron feature size, 2250 transistors
-vs. eg, **2020** Intel Xeon W-3175X has .014 micron sizes,
8,000,000,000 (8B) transistors, 64bit, 4300Mhz (4GHz)clock

MAKING CHIPS – NOT MAGIC, JUST ADVANCED PHOTOGRAPHY




160

October 2020

- START WITH MEMORY LOCATION 0
- 1'S AND 0'S INTO LOGIC CIRCUITS IN CPU
- CPU -- DO SOME SIMPLE OPERATION

EG, memory A > memory B? Yes, write a '1'


EG, Jump to memory location xx



SOME OF THE 1'S & 0'S ARE DATA, SOME OF THEM SPECIFY LOGIC OPERATIONS THE CPU IS SUPPOSED TO DO

EG, COMPARE TWO MEMORY LOCATIONS ALREADY LOADED IN THE CPU (EG, LOCATIONS 58 & 96)

-IF LOCATION 58 > LOCATION 96
THEN SET A MEMORY 'FLAG'
WHICH THEN NEXT STEP CAUSES
A JUMP TO MEMORY LOCATION
XX AND USE THIS AS THE NEXT
LOGIC OPERATION




SYMBOLIC COMPUTER PROGRAM ('GOFAI') IS TURNED INTO 1'S AND 0'S AND CPU PROCESSES IT

King(John) \wedge Greedy(John) \Rightarrow Evil(John)

```
if king and greedy:  
    print(f'{x} is an evil king')  
else:  
    .....  
.....
```

PYTHON
COMPILER/
INTERPRETER




NEW TOPIC

SECOND REVIEW QUIZ

164

October 2020



QUIZ – QUESTION 1.

WHAT IS REINFORCEMENT LEARNING?

166

October 2020


REINFORCEMENT LEARNING

-NETWORK ('AGENT') TRIES TO MAXIMIZE THE REWARD
EG, PLAYS CHESS/SHOGI/GO AGAINST ITSELF MILLIONS
AND MILLIONS OF TIMES




Deepmind 2017: Alphazero –
superhuman performance in
chess, shogi & go

QUIZ – QUESTION 2. WHAT KIND OF SYSTEM IS SHOWN BELOW?



DEEP LEARNING TO RECOGNIZE ALL SORTS OF CATS AND DOGS



“Connectionism”
“Neural Network”
“Deep Learning Network”

DEEP LEARNING CAN RECOGNIZE OR ASSOCIATE ALL KINDS OF INFORMATION – IMAGES, FACES, SPEECH, MARKETING DATA, XRAY IMAGES.....



QUIZ – QUESTION 3. WHAT KIND OF COMPUTER SYSTEM HANDLES THE LOGIC SHOWN BELOW?

$\forall x \text{King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$ *← axiom*
(“All greedy kings are evil”)

Therefore can infer:

$\text{King}(\text{John}) \wedge \text{Greedy}(\text{John}) \Rightarrow \text{Evil}(\text{John})$

GOOD OLD FASHIONED ARTIFICIAL INTELLIGENCE ("GOFAI") "SYMBOLIC AI"

$\forall x \text{King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$ *← axiom*
("All greedy kings are evil")

Therefore can infer:

$\text{King}(\text{John}) \wedge \text{Greedy}(\text{John}) \Rightarrow \text{Evil}(\text{John})$

172


GOOD OLD FASHIONED ARTIFICIAL INTELLIGENCE “GOFAI” == “SYMBOLIC AI”

```
yes_synonyms = ['y', 'Y', 'yes', 'Yes', 'T', 'true']
x = input('What is the name of the person? ')
king = input(f'Is {x} a king? ')
if king in yes_synonyms:
    king = True
else:
    king = False
greedy = input(f'Is {x} greedy? ')
if greedy in yes_synonyms:
    greedy = True
else:
    greedy = False
if king and greedy:
    print(f'{x} is an evil king')
else:
    print(f'We cannot infer that {x} is an evil king')
```

```
C:\Users\howar>kings.py
What is the name of the
person? Sebastian
Is Sebastian a king? yes
Is Sebastian greedy? yes
Sebastian is an evil king
```


Quiz – question 4

Where are Neural Networks in the diagram below?
Where is Symbolic (GOFAI) in the diagram below?



Neural networks

Symbolic AI



LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ **1c. Neuro-Symbolic Gap**
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. How will AI affect my practice of medicine including psychotherapy?
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. How will AI affect the future of mankind?
- ▶ 5. Discussion

NEW TOPIC

NEURO-SYMBOLIC GAP

177

October 2020

WE HAVE COVERED :

-DEEP LEARNING (**NEURAL** NETWORKS AI)
ASSOCIATIONS, RECOGNITION (IMAGES, SPEECH, ETC....)
BIG SUCCESSES!!

-**SYMBOLIC** AI
LOGIC

BIG FAILURES!! EG, EXPERT SYSTEMS -> AI WINTERS!!

178

October 2020

The Neural Symbolic Gap




- **Neural Network** – phenomenal image processing and reinforcement learning
- **Child** – phenomenal causal symbolic learning with few examples (eg, Gopnik)



“panda”
57.7% confidence

+



=




“gibbon”
99.3 % confidence

Goodfellow,I.J., Shlens,J. and Szegedy,C. (Google Mountainview), Explaining and Harnessing Adversarial Examples, ICLR 2015.

It's still a Panda – and the 3 year old boy would know this!!
(and.... 3 year old only needs 1 or 2 photos for training, not 1000s)

Dr Howard Schneider - AI & Future of Medicine **INTENDED FOR LIVE PRESENTATION, NOT FOR SELF-READING**





Deep Learning Neural Network

Pattern Recognition
→Recognize the World

Need 1000's examples for learning

3 Year Old Human Child

Model Building +also Pattern Recognition
→Explain the World

A few examples enough


181

Chimpanzee (behind plexiglass)

Stick to push or pull food

Food

Gravity trap



Chimpanzees **do not** have full causality

182

October 2020

Deep Learning has zero Understanding
**Deep Learning gets more powerful by
brute force**

300,000x increase in computing power
the last few years
NOT SUSTAINABLE -- There is a deep flaw
in deep learning

183

October 2020

Humans have causality. ✓

Animals do not. ← associations

Deep Learning does not. ← associations

184

October 2020

Humans learn to drive with 5-15 hours of driving in a driver's ed course. Do not need to be programmed with every possible scene on the road – can figure it out (eg, bag on road).

Deep Learning autonomous driving system – must get training data showing *everything* since **it cannot causally reason**
eg, one EV company has about a billion miles of driving experience of its cars as training data now.


185

Classic example: unless explicitly programmed, an autonomous driving system will follow a seemingly normal road and drive off a cliff (because it understands nothing!!)

(programmers add symbolic rules to autonomous driving systems and probably this one has been added already)



Without causality need loads and loads of experience to see everything: ‘long-tail problem’



March 18, 2018 – Uber self driving car kills pedestrian in Tempe, Arizona

-Uber's system never had training with image of pedestrian pushing bike across the *middle* of the street
(Safety driver did not unfortunately respond to situation)



The solution to the neurosymbolic gap: Ability to Generate Causal Behavior



‘Reptilian’ and ‘Mammalian’ Brain
– Associative, Pre-Causal
Functioning




‘Human’ Brain, AGI – Full
Causal Functioning

Directed Acyclic Graphs ('Causal Graphs')

Counterfactual Theory


- Useful for **Analyzing** Causality, eg, epidemiology, genomics
- Less Useful for **Generating** Causality, eg, AGI



Causal Discovery

Try to learn causal relationships from the data

eg, Data shows A & B are independent if we condition on C,
but dependent if we don't



other causal models will also exist given this data....

Directed Acyclic Graphs ('Causal Graphs')

Counterfactual Theory

Causal Discovery


-Useful for **Analyzing** Causality

eg, epidemiology, genomics

-Less Useful for **Generating** Causal Behavior


eg, AGI

We want a mechanism for generating causal behavior in the real world



CAUSAL COGNITIVE ARCHITECTURE 1 – SCHNEIDER (IN PRESS)

Connectionist elements
Integrated into a navigation-based framework. Causality emerges from the architecture.



LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. **How will AI in next decade (or two) affect my patients' lives?**
- ▶ 3. **How will AI affect my practice of medicine including psychotherapy?**
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. **How will AI affect the future of mankind?**
- ▶ 5. **Discussion**

NEW TOPIC

HOW WILL AI IN THE FUTURE AFFECT PATIENTS....ACTUALLY AFFECT ALL OF US?

196

October 2020

THIS DECADE (MY OPINION)

- OTHER FACTORS OF MORE IMPORTANCE
EG, COVID-19 RECOVERY, GLOBALIZATION, ETC**
- IMPROVEMENTS IN OFFICE AUTOMATION AND
FACTORY AUTOMATION (EVEN WITHOUT AI) WILL
AFFECT JOBS (IN GOOD AND BAD WAYS)**
- DEEP LEARNING CAN SPEED THIS UP**

197

October 2020



AI AND THE WORK OF
THE FUTURE CONGRESS
2019
UNCONFERENCE
REPORT

"SENSE OF POWERLESSNESS – WORRY..."
→ RECORD LEVELS OF INEQUALITY, LOWER SOCIAL
MOBILITY

198

October 2020

**-HOWEVER, TECH USE IN HIGH SKILL JOBS
MAKES THESE JOBS MORE PRODUCTIVE &
PAYING VS LOW SKILL JOBS**

**→NEED TO FIX IMBALANCE BETWEEN
WORKERS HELPED AND HURT BY NEW
ECONOMY**

→ NOT REALLY DUE TO AI AT THIS POINT

199

AI -- NEXT DECADE


*“It's tough to make predictions,
especially about the future”*


--Yogi Berra

200

October 2020

Gartner Hype Cycle for Artificial Intelligence, 2019





NEXT DECADE (MY OPINION)

-ONCE THERE IS CAUSALITY IN AI + APPROPRIATE TRAINING + LOWERED 'COMPUTE' COSTS -- AI SYSTEMS SHOULD BE ABLE TO DO THE WORK ALMOST ANY WORKER CAN DO

-TIME OF STRESS FOR OUR PATIENTS WITHOUT APPROPRIATE SOCIETY PROGRAMS IN PLACE

203

October 2020

LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. **How will AI affect my practice of medicine including psychotherapy?**
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. **How will AI affect the future of mankind?**
- ▶ 5. **Discussion**

LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. How will AI affect my practice of medicine including psychotherapy?
- ▶ **3a. How is AI affecting medicine at present?**
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. How will AI affect the future of mankind?
- ▶ 5. Discussion

205

October 2020

NEW TOPIC

HOW WILL AI IN FUTURE AFFECT THE PRACTICE OF MEDICINE, INCLUDING PSYCHOTHERAPY?

206

October 2020

MEDICAL AI SYSTEMS


- NO UNDERSTANDING OF WHAT THEY ARE DOING
- NO/POOR CAUSALITY
- THEY CAN'T EVEN EXPLAIN WHAT THEY ARE DOING



207

October 2020

COVER STORY OF JANUARY 2020 PRINT CMAJ



Dr Howard Schneider - AI & Future of

208

October 2020

THE ARTICLE TALKED ABOUT INTELLIGENCE IN HUMANS AND MACHINES IN VAGUE WAYS....

ANALYSIS+

“Because human and artificial intelligences are different and complementary, it is unlikely that AI will entirely replace the physician in the resolution of clinical problems.”



My view:

- current AI contains **no model** of the world, **no causal reasoning**
- instead functions as a **massive association machine**

DEEP LEARNING FOR RADIOLOGY (REQUIRES MILLIONS OF XRAYS AS TRAINING DATA)

- USEFUL 3AM IN ER
 - USEFUL SCREENING
 - BUT.... NO CAUSALITY
- CLINICAL CORRELATION
DISEASE PROGRESSION
TEAM COLLABORATION

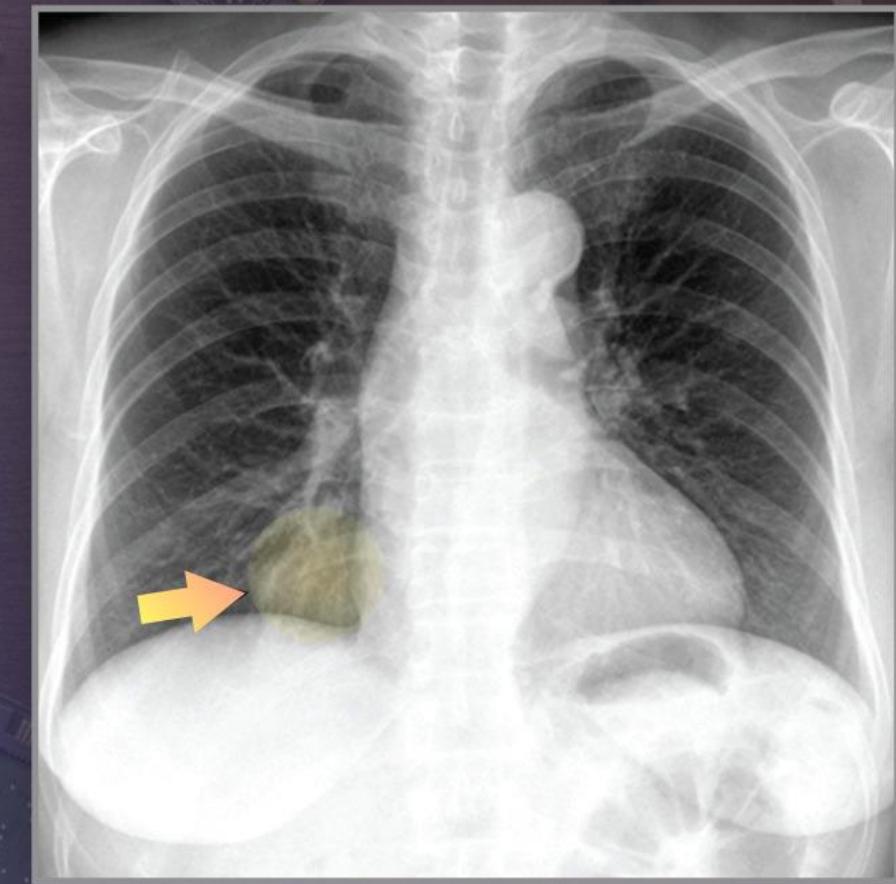
Human vs Machine: Lung Tumor

Chest X-Rays image the lungs, heart, blood vessels, and bones. AI has been used to read and understand them.

Example:
Lung Tumor

Computers:
Score: 0.291

Doctors:
2/15 Detected



Artificial Intelligence in Medicine

Volume 103, March 2020, 101785

The impact of machine learning on patient care: A systematic review


Ben-Israel and colleagues - Neurosurgery, University of Calgary

“Despite the expanding use of machine learning (ML) in fields such as finance and marketing, its application in the daily practice of clinical medicine is almost non-existent.”

211


October 2020

Dr Sidney Kennedy, University of Toronto
Machine Learning in psychiatry
-predict escitalopram treatment outcome from EEG
-predict relapse, increased risk of suicidal behavior



AI Therapists

- OFFERS PROMISE OF THERAPY FOR ALL AT LOW COST
- ONLINE CBT
- CHATBOTS: "SIRI, I FEEL DEPRESSED"
- CHATBOTS, EG, WOEBOT



Computer-Assisted CBT ('CCBT') for Depression

- JESSE WRIGHT ET AL, 2019, J CLIN PSYCHIATRY
- META-ANALYSIS 40 RCT STUDIES CCBT FOR DEPRESSION
- CCBT WITH MODEST SUPPORT FROM CLINICIAN EFFECTIVE
- CCBT FULLY SELF-GUIDED MUCH LESS EFFECTIVE
(TYPES, AMOUNT AI IN CCBT STUDIED??)

214

October 2020

→ AT THIS POINT IN TIME AI CANNOT
REPLACE DOCTORS -- NOT TECHNICALLY
POSSIBLE

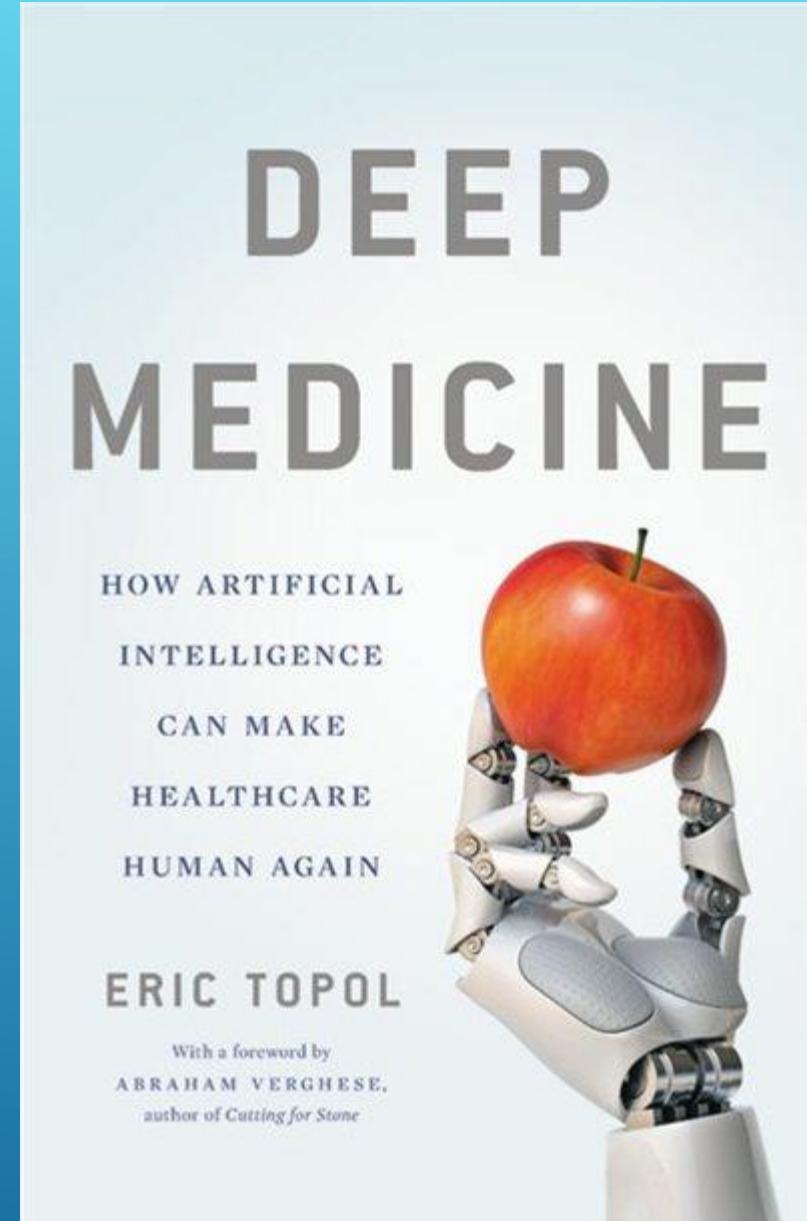
- BETTER SOFTWARE CAN IN THEORY REDUCE
HEALTH CARE STAFF HOWEVER, AND DEEP
LEARNING CAN SPEED THIS UP (DEPENDS ON
FIELD)

215

October 2020

DR ERIC TOPOL

-CARDIOLOGIST
-SCRIPPS, LA JOLLA
-DEVELOPED T-PA,
ANTI-CLOT MEDS
**-PROPOONENT OF
AI IN MEDICINE**



LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. How will AI affect my practice of medicine including psychotherapy?
- ▶ 3a. How is AI affecting medicine at present?
- ▶ **3b. How will AI affect medicine in the next decade?**
- ▶ 4. How will AI affect the future of mankind?
- ▶ 5. Discussion

NEXT DECADE (MY OPINION)

- ONCE THERE IS CAUSALITY + APPROPRIATE FUNDAMENTAL TRAINING + REDUCED COSTS FOR COMPUTE -- AI SYSTEMS SHOULD BE ABLE TO DO THE WORK ALMOST ANY HEALTHCARE WORKER CAN DO
- HUMANS PROVIDE **OVERSIGHT** (FOR A WHILE....) OF AI HEALTH SYSTEMS

218

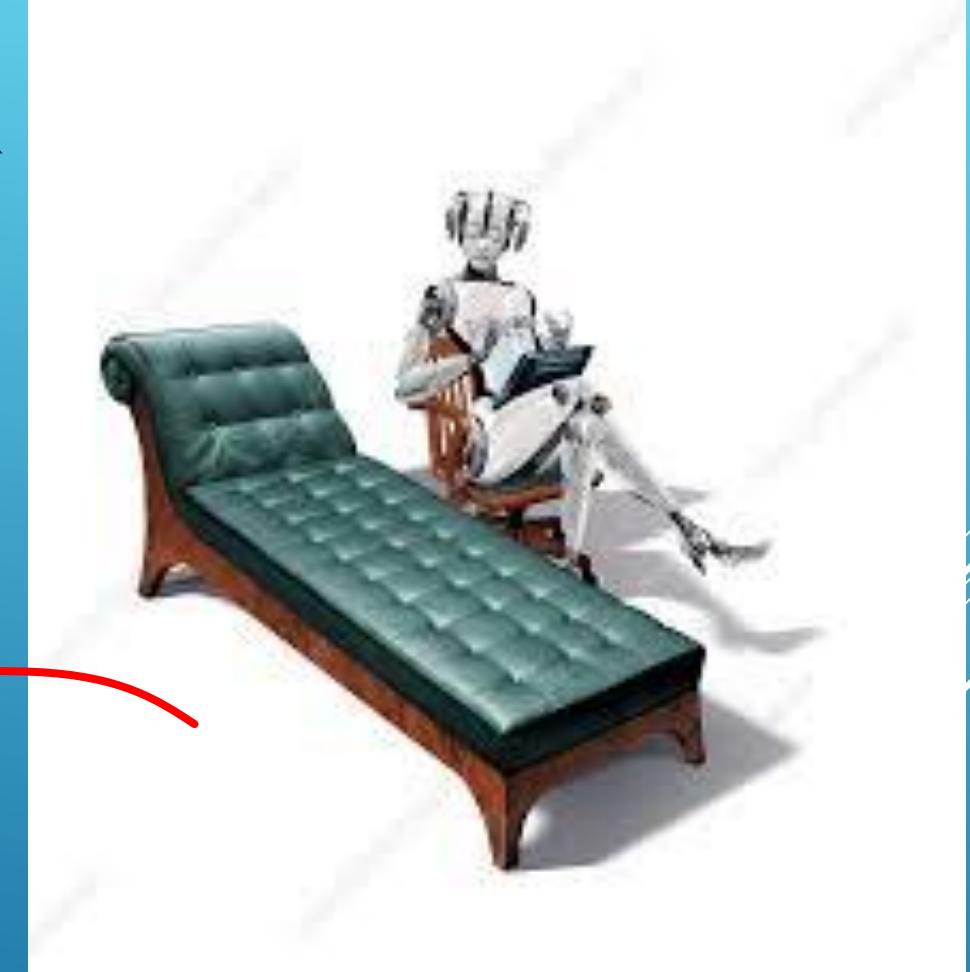
October 2020

NEXT DECADE (MY OPINION)

HUMANS PROVIDE EMPATHIC &
EXISTENTIAL LIAISON TO AI
HEALTH SYSTEMS

→ THUS KEY FOR
PSYCHOTHERAPY

Even if technology is
great, will this work?



NEW TOPIC

HOW WILL AI AFFECT THE FUTURE OF MANKIND (HUMANKIND)?

220

October 2020

LEARNING OBJECTIVES

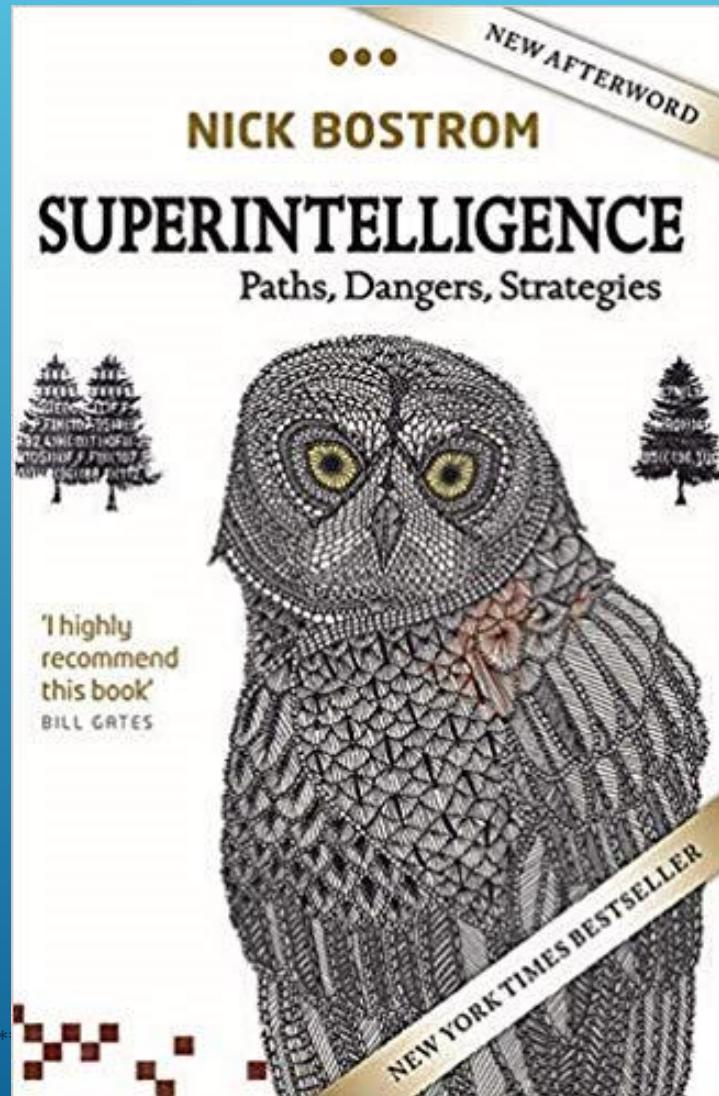
- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. How will AI affect my practice of medicine including psychotherapy?
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. How will AI affect the future of mankind?
- ▶ 5. Discussion

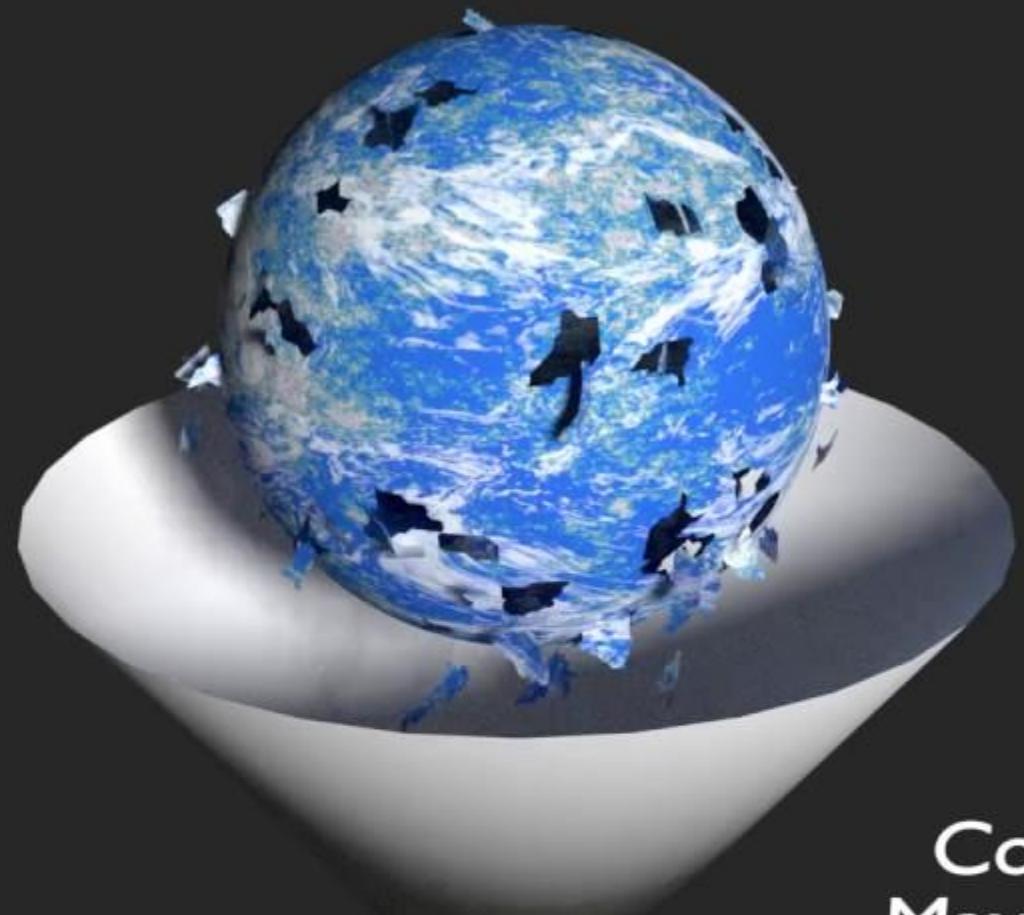
DEEP LEARNING ALL AROUND US NOW: LONDON POLICE SURVEILLANCE CAMERAS WITH FACIAL RECOGNITION

PSYCHOLOGY OF CONSTANT SURVEILLANCE??



Bostrom – SuperIntelligence






*Could an AI Paperclip
Maximizer machine turn
the earth into paperclips?*

224

October 2020

Kurzweil – AI Will Exceed Human Intelligence and Merge with Humans



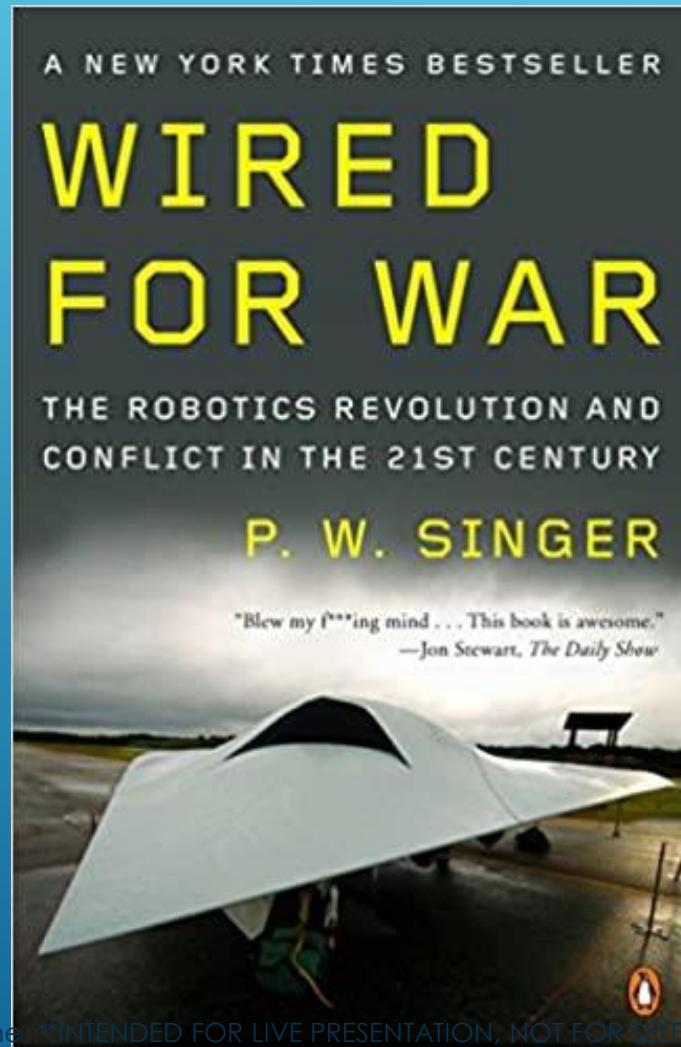
Dr Howard Schneider - AI & Future of Medicine **INTL

225

October 2020



SINGER'S 2009 BOOK – WIRED FOR WAR



Dr Howard Schneider - AI & Future of Medicine **INTENDED FOR LIVE PRESENTATION, NOT FOR SELF-READING**

2017 ASILOMAR AI PRINCIPLES

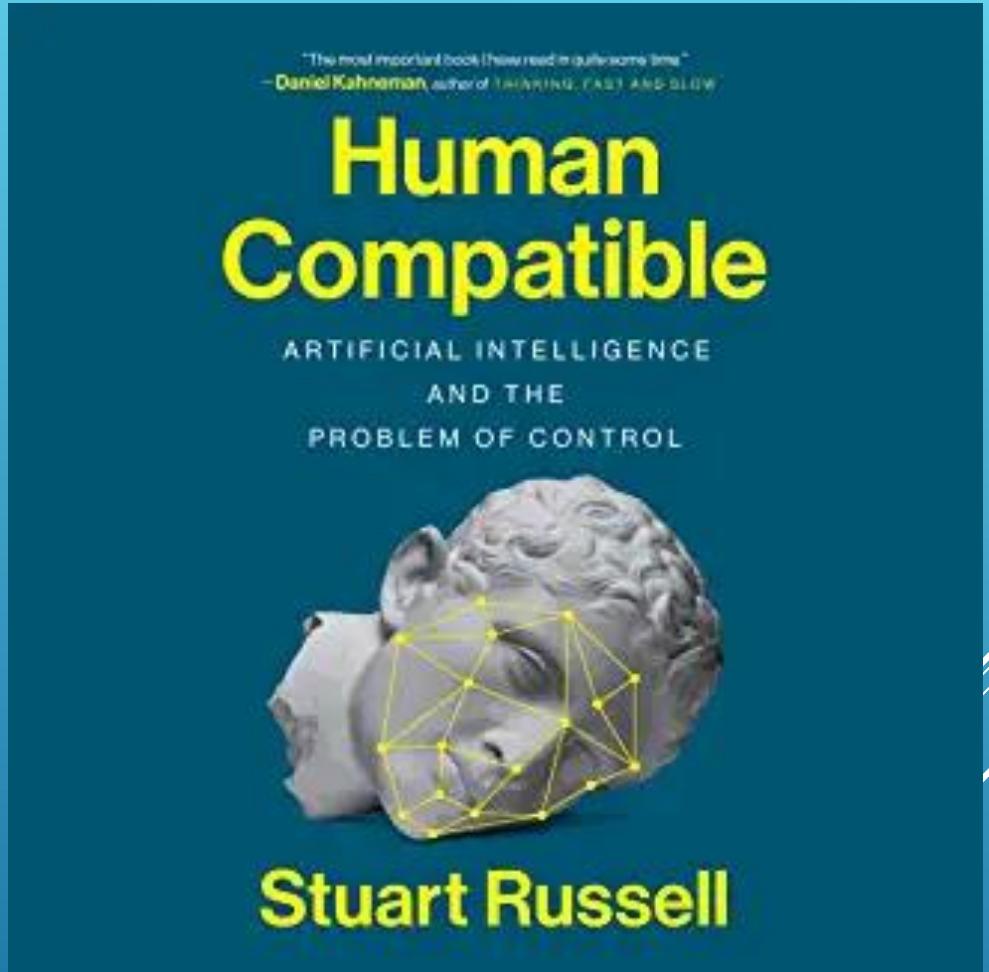


- #1- The goal of AI research should be to create not undirected intelligence, but **beneficial intelligence**.
- #6- AI systems should be **safe and secure.. verifiably so..**
- #23- **Superintelligence** should only be developed.. for the **benefit of all humanity** rather than one state.


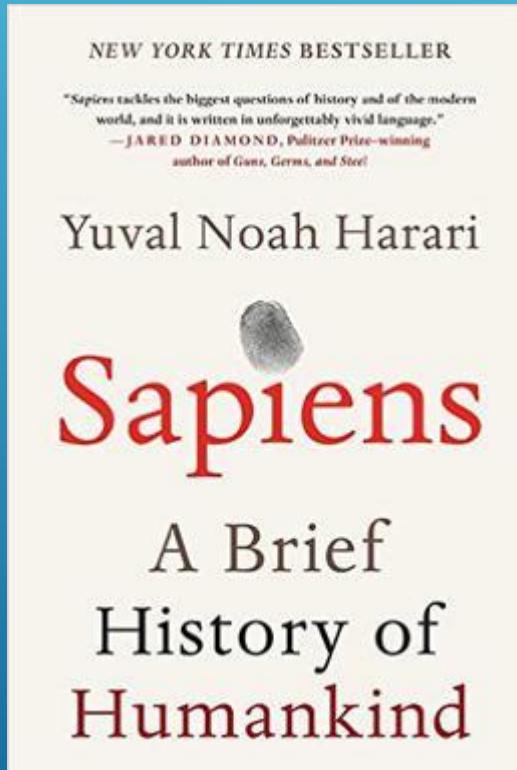
228

Provably Beneficial AI

The only objective
of the machine is
human well being,
or the satisfaction
of human
preferences



- Continued technological powers.
- We lose meaning in our lives
- Replacement of *H sapiens* with *H deus*



RIGHT NOW

- AI (= DEEP LEARNING) IS A PHENOMENAL COMPUTER SCIENCE TECHNIQUE
- LIKELY TO **IMPROVE OUR LIVES**
- TERMINATOR NOT POSSIBLE

FUTURE

- THAT'S ANOTHER STORY
- **AI SAFETY CONCERNS**

OK.... WE'VE MADE IT....



NEW TOPIC

DISCUSSION

233

October 2020

LEARNING OBJECTIVES

- ▶ 1. Real understanding of what AI is:
- ▶ 1a. Deep Learning and Reinforcement Learning
- ▶ 1b. Field of Artificial Intelligence (AI)
- ▶ 1c. Neuro-Symbolic Gap
- ▶ 2. How will AI in next decade (or two) affect my patients' lives?
- ▶ 3. How will AI affect my practice of medicine including psychotherapy?
- ▶ 3a. How is AI affecting medicine at present?
- ▶ 3b. How will AI affect medicine in the next decade?
- ▶ 4. How will AI affect the future of mankind?
- ▶ 5. Discussion

Artificial Intelligence and the Future of Medicine/Mankind

Thank You

howard.schneider@gmail.com

To download a copy of these slides:

SEARCH (EG, GOOGLE):

“DR HOWARD SCHNEIDER TORONTO”

OR ENTER:

[HTTPS://GITHUB.COM/HOWARD8888/AI-AND-THE-FUTURE-OF-MEDICINE/BLOB/MASTER/AI%20AND%20THE%20FUTURE%20OF%20MEDICINE.PDF](https://github.com/howard8888/AI-AND-THE-FUTURE-OF-MEDICINE/blob/master/AI%20AND%20THE%20FUTURE%20OF%20MEDICINE.PDF)



236